



Motion Planning

Nicolas Perrin-Gilbert

One of the reference books:

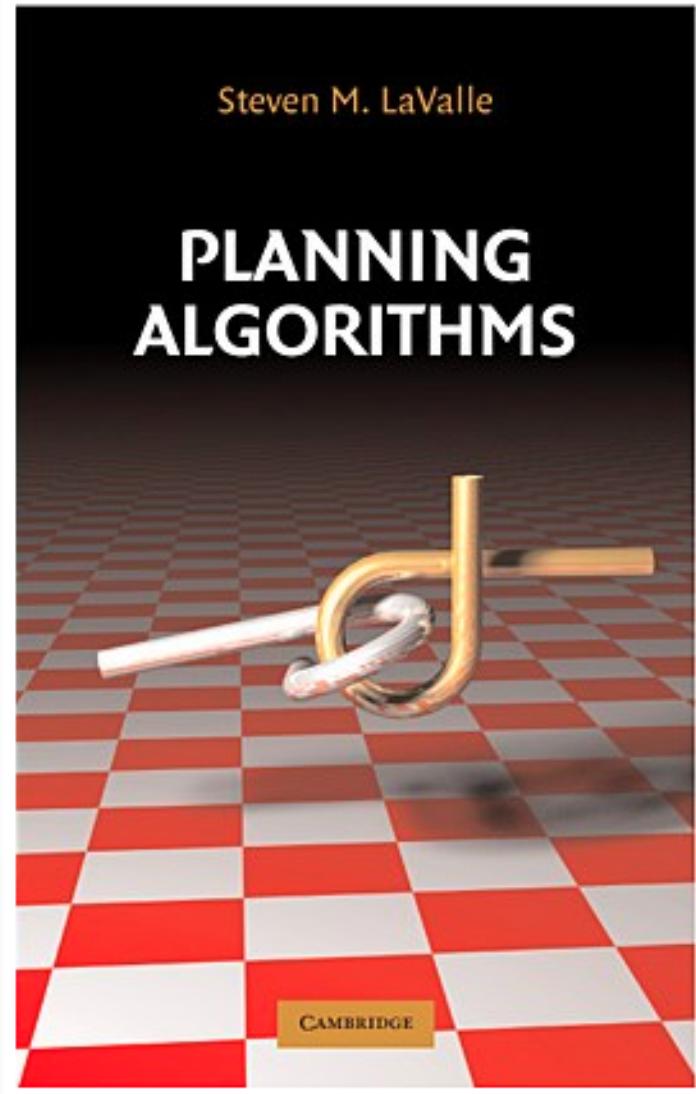
Planning Algorithms

Steven M. LaValle

Cambridge University Press

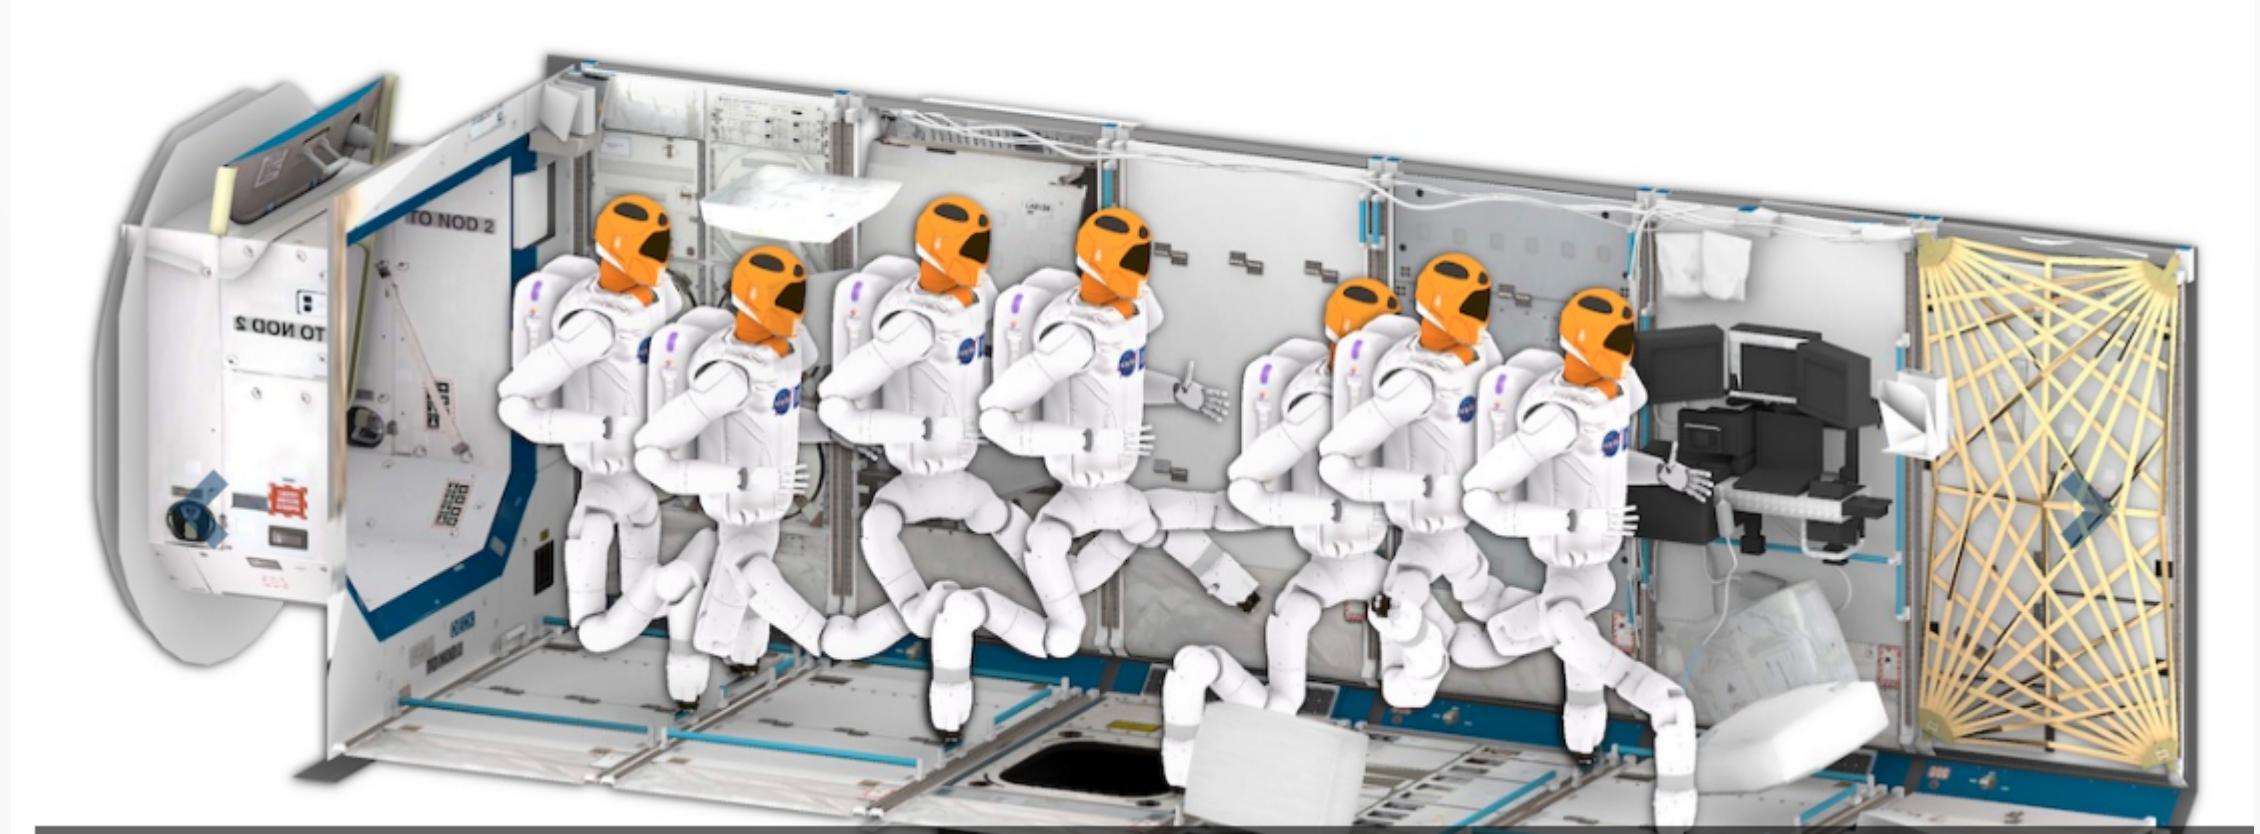
Available at:

<http://lavalle.pl/planning/>



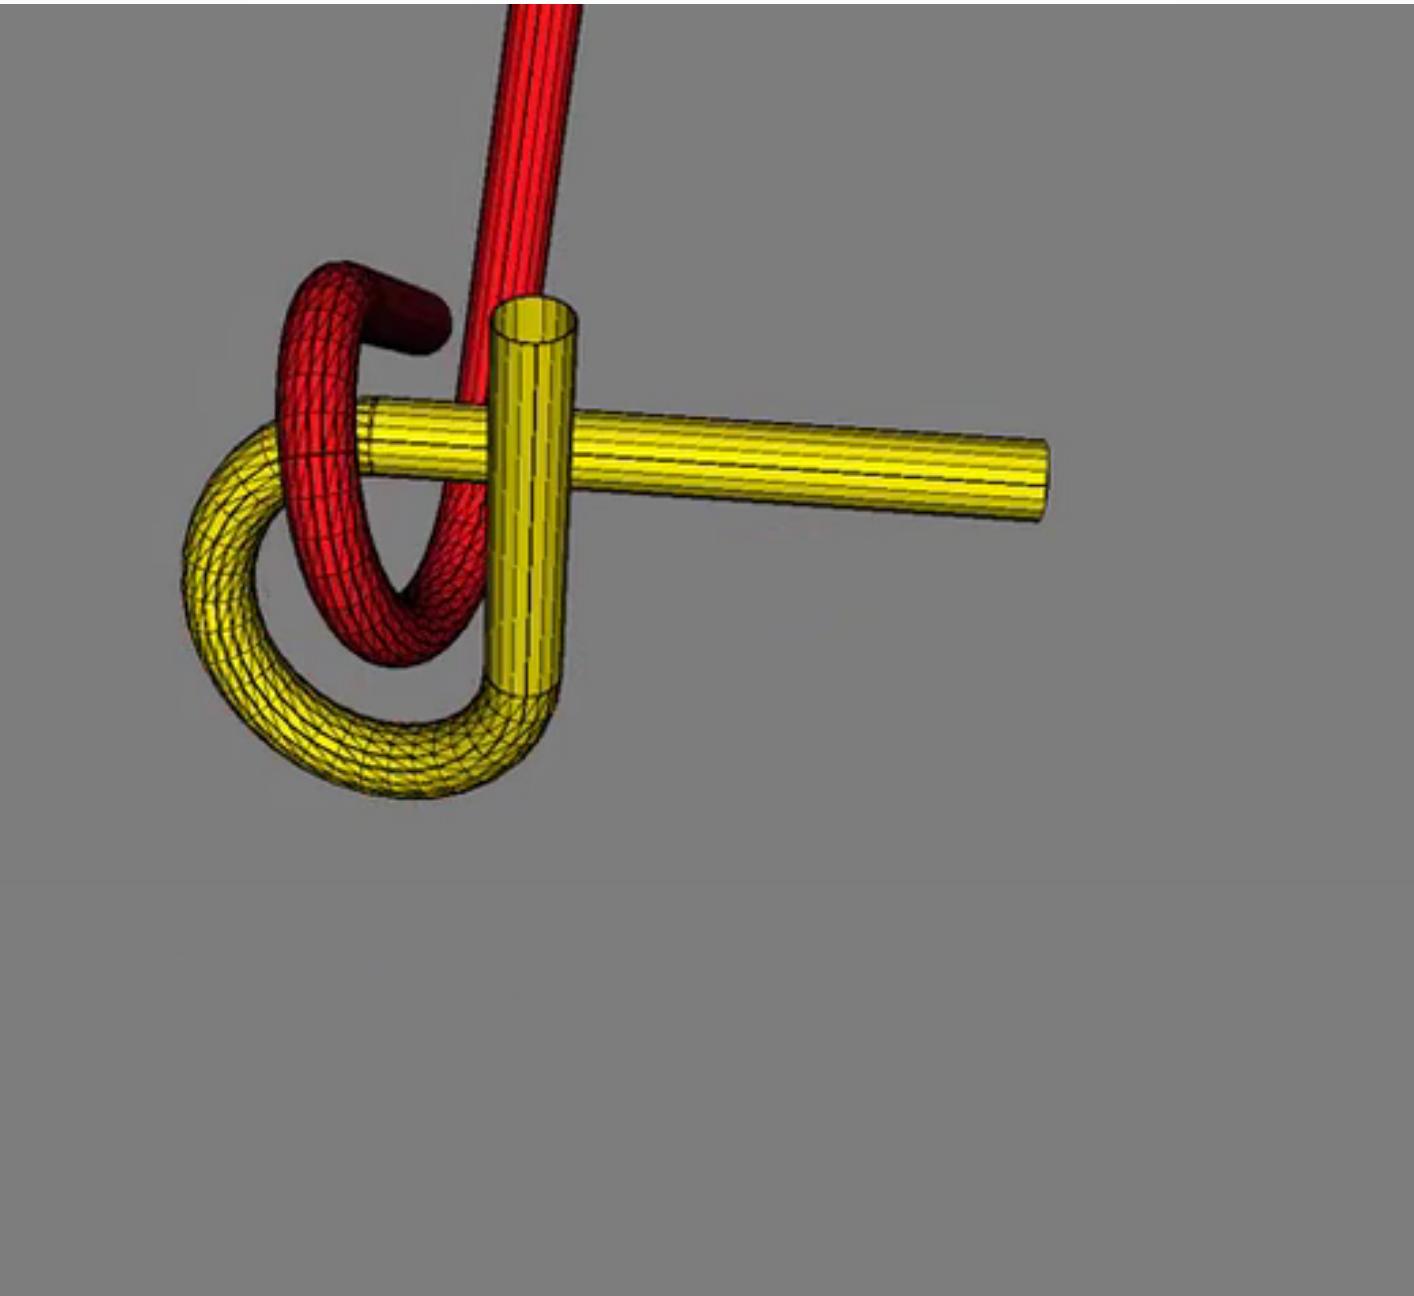
Reference library (C++/Python):

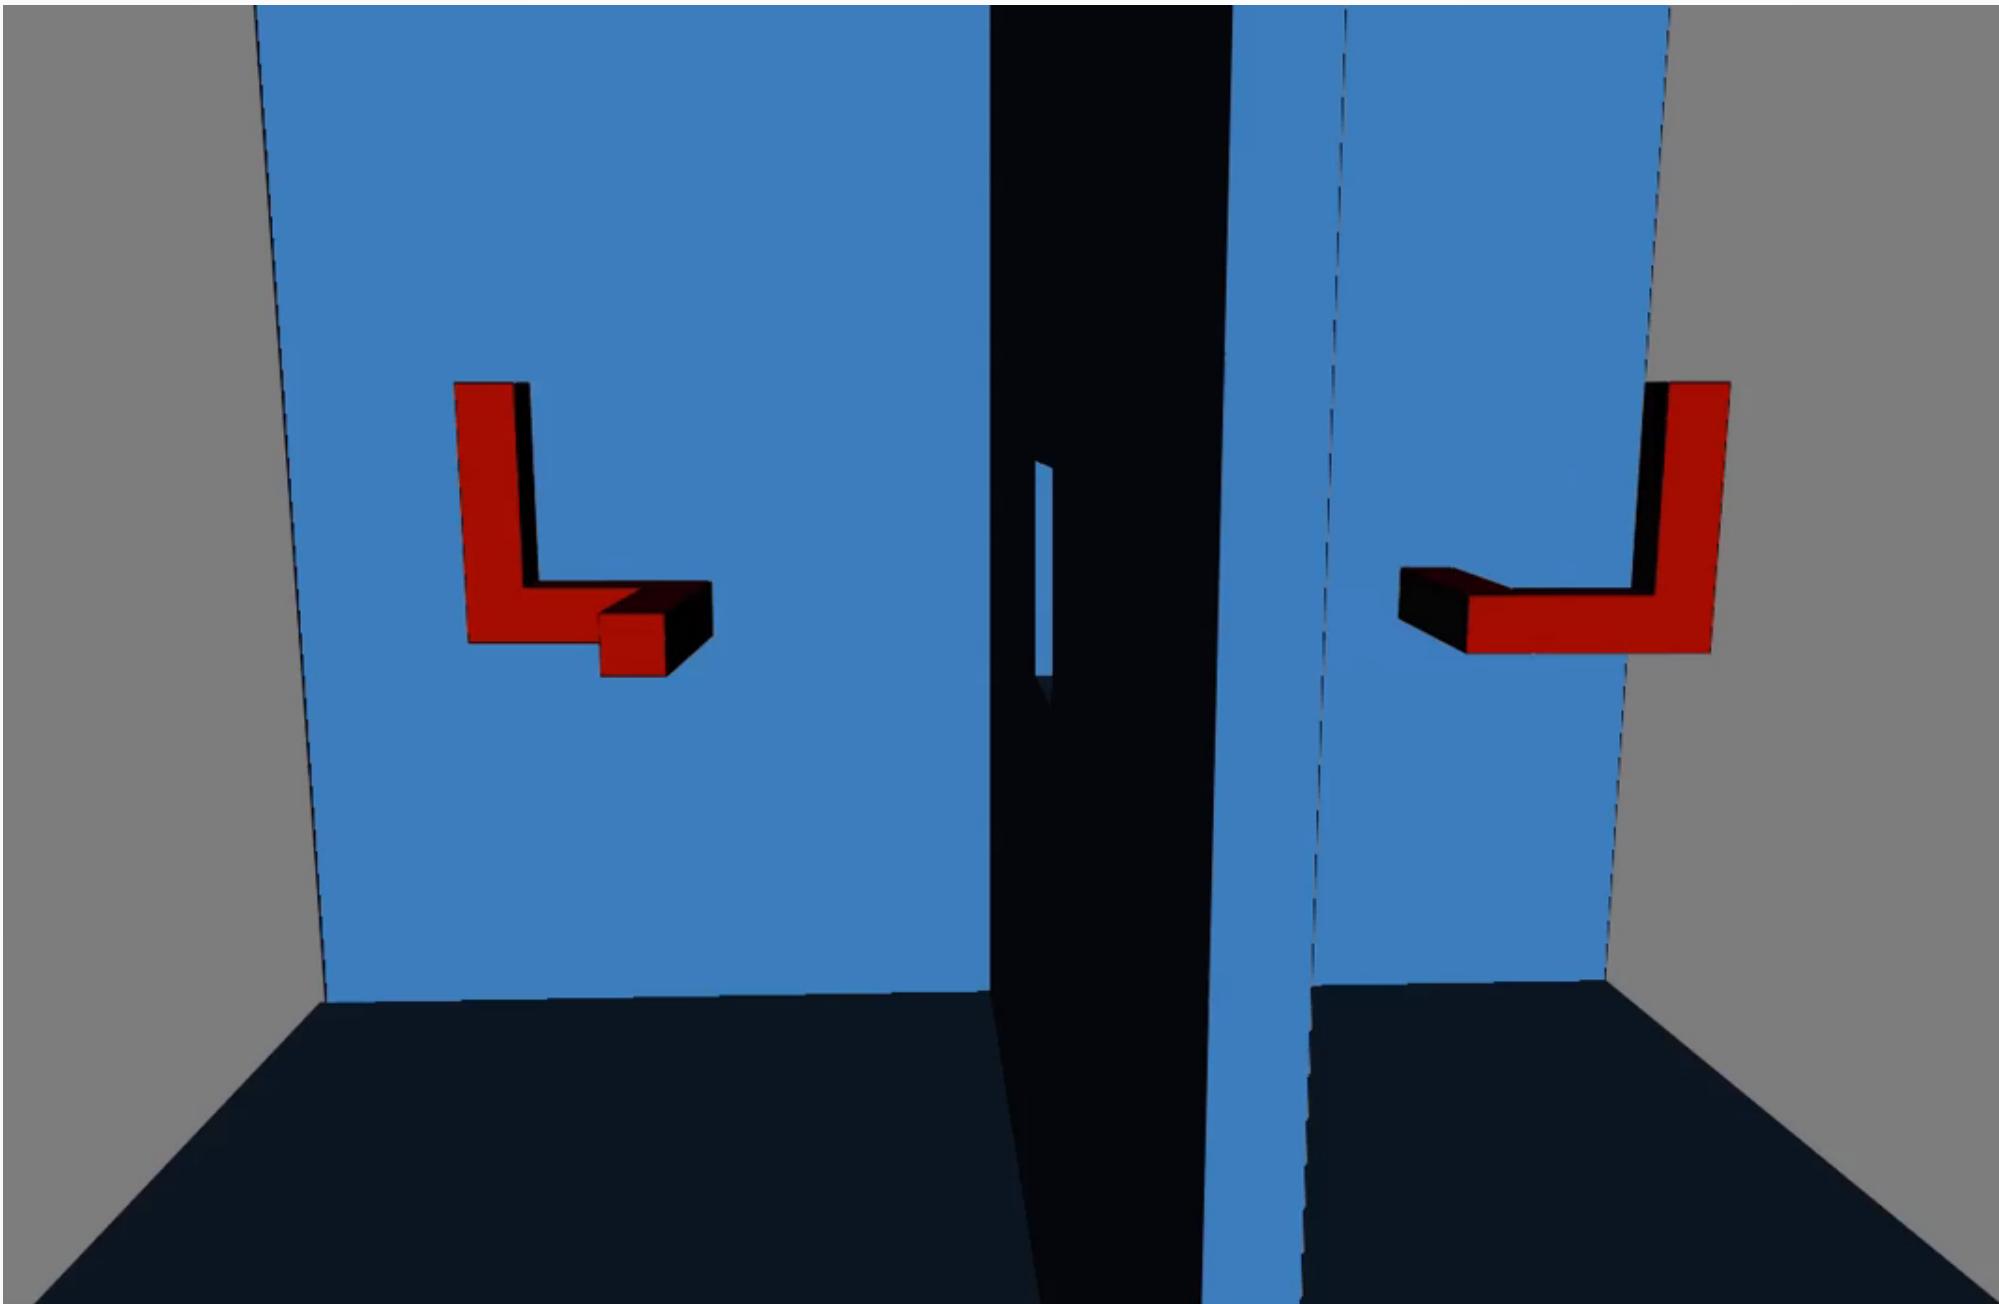
The Open Motion Planning Library

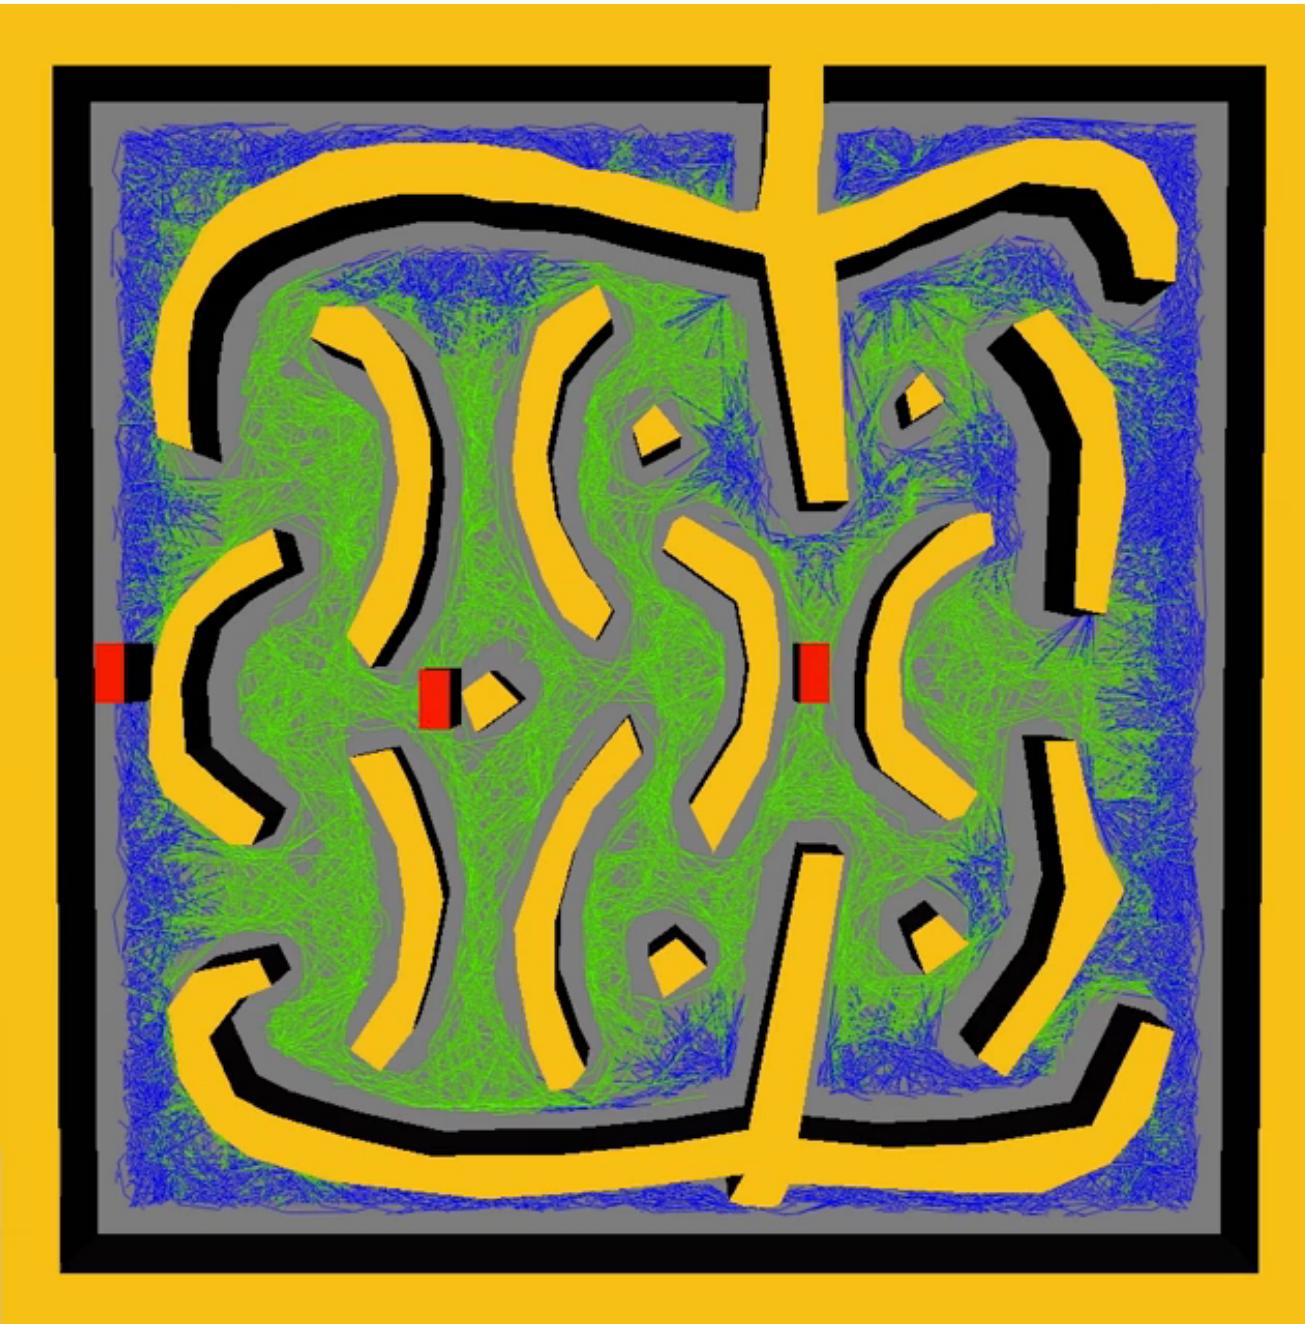


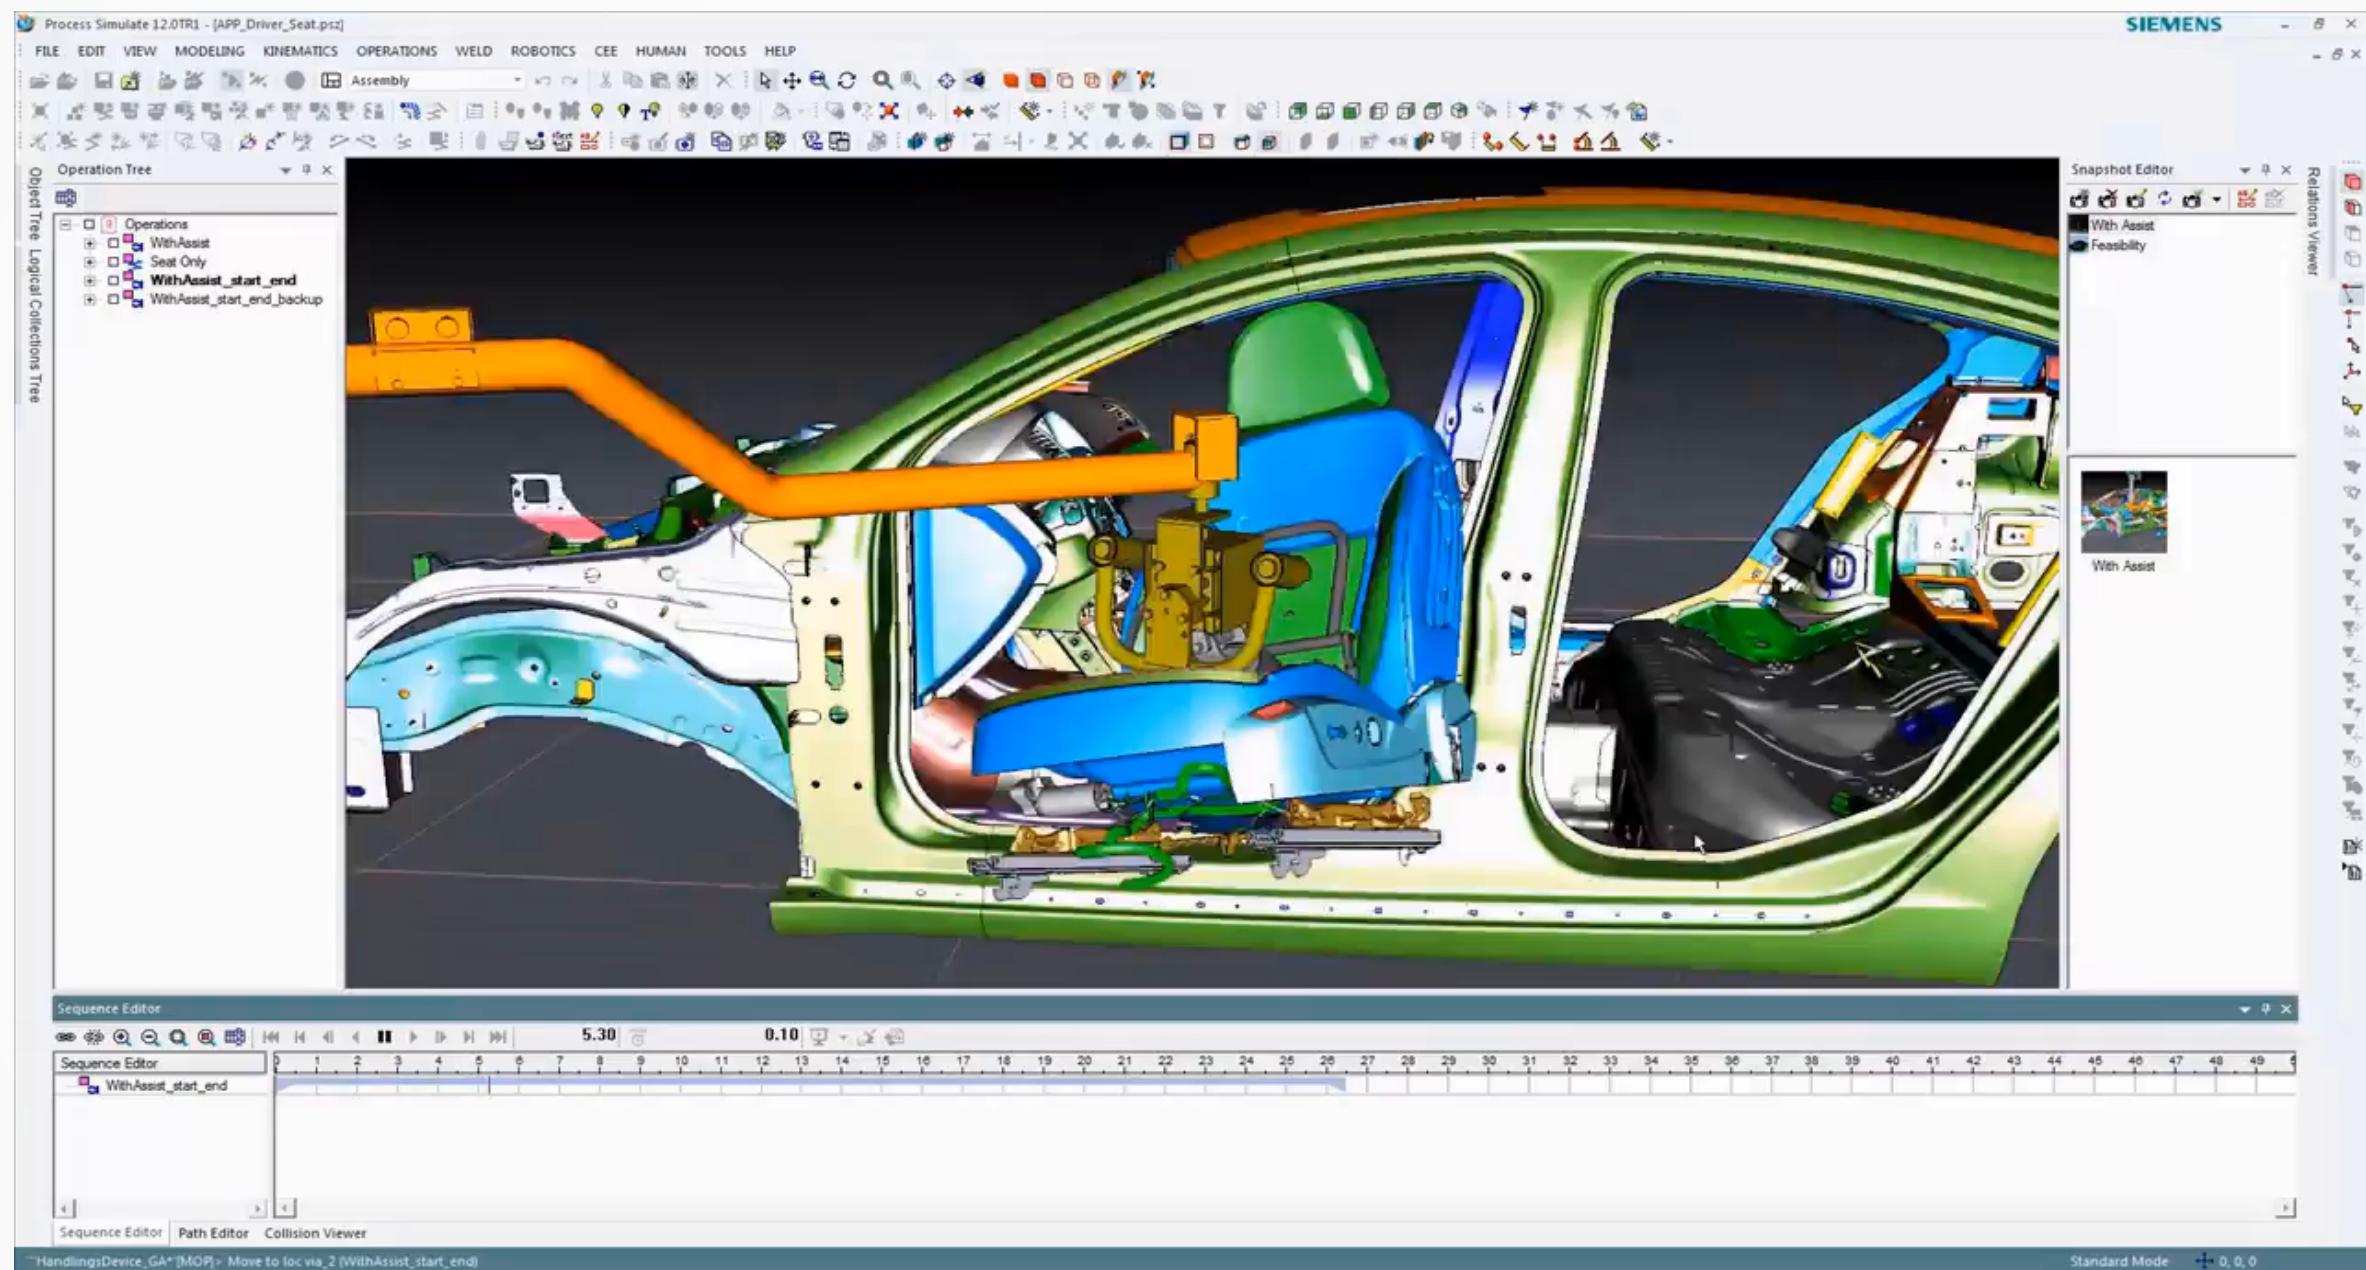
Website: **ompl.kavrakilab.org**









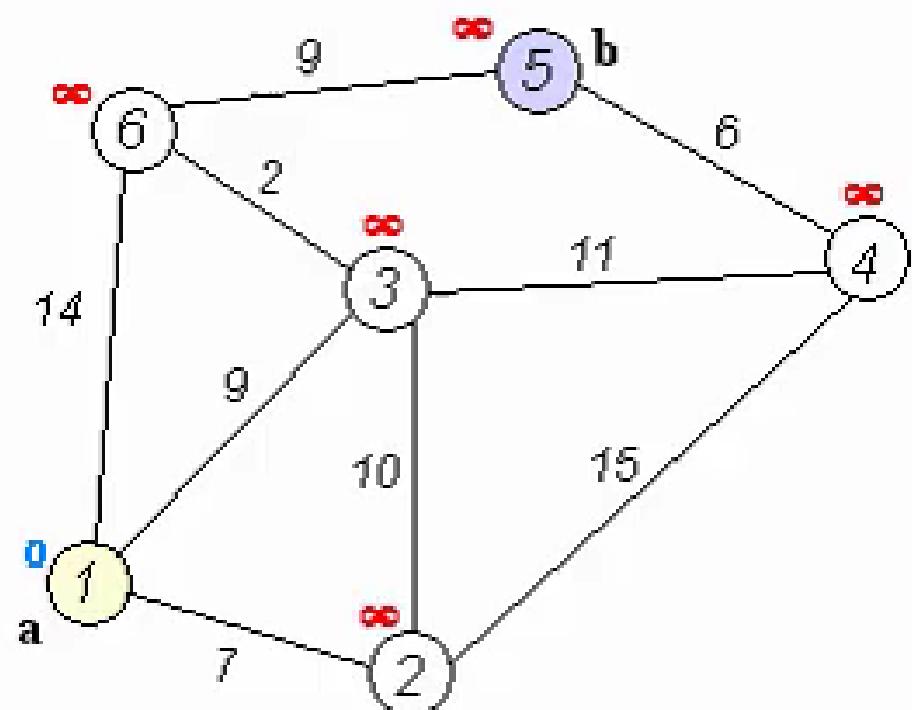


Preamble: path planning in a graph

Shortest path: **Dijkstra's algorithm** (1959)

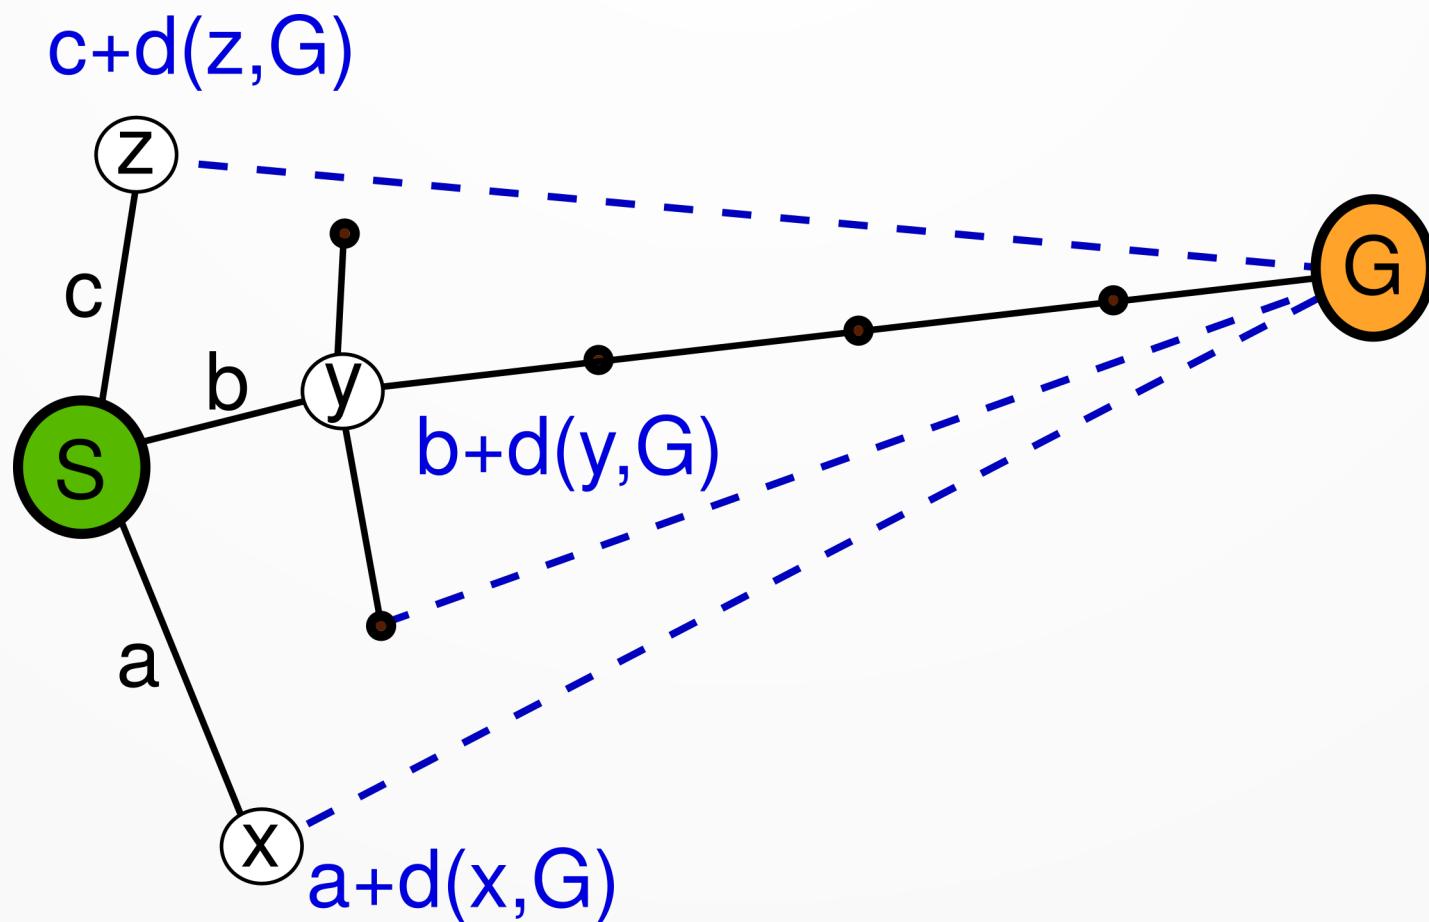
Complexity: $O(a+n \log n)$

→ Slow if there are many vertices, or many edges.



Preamble: path planning in a graph

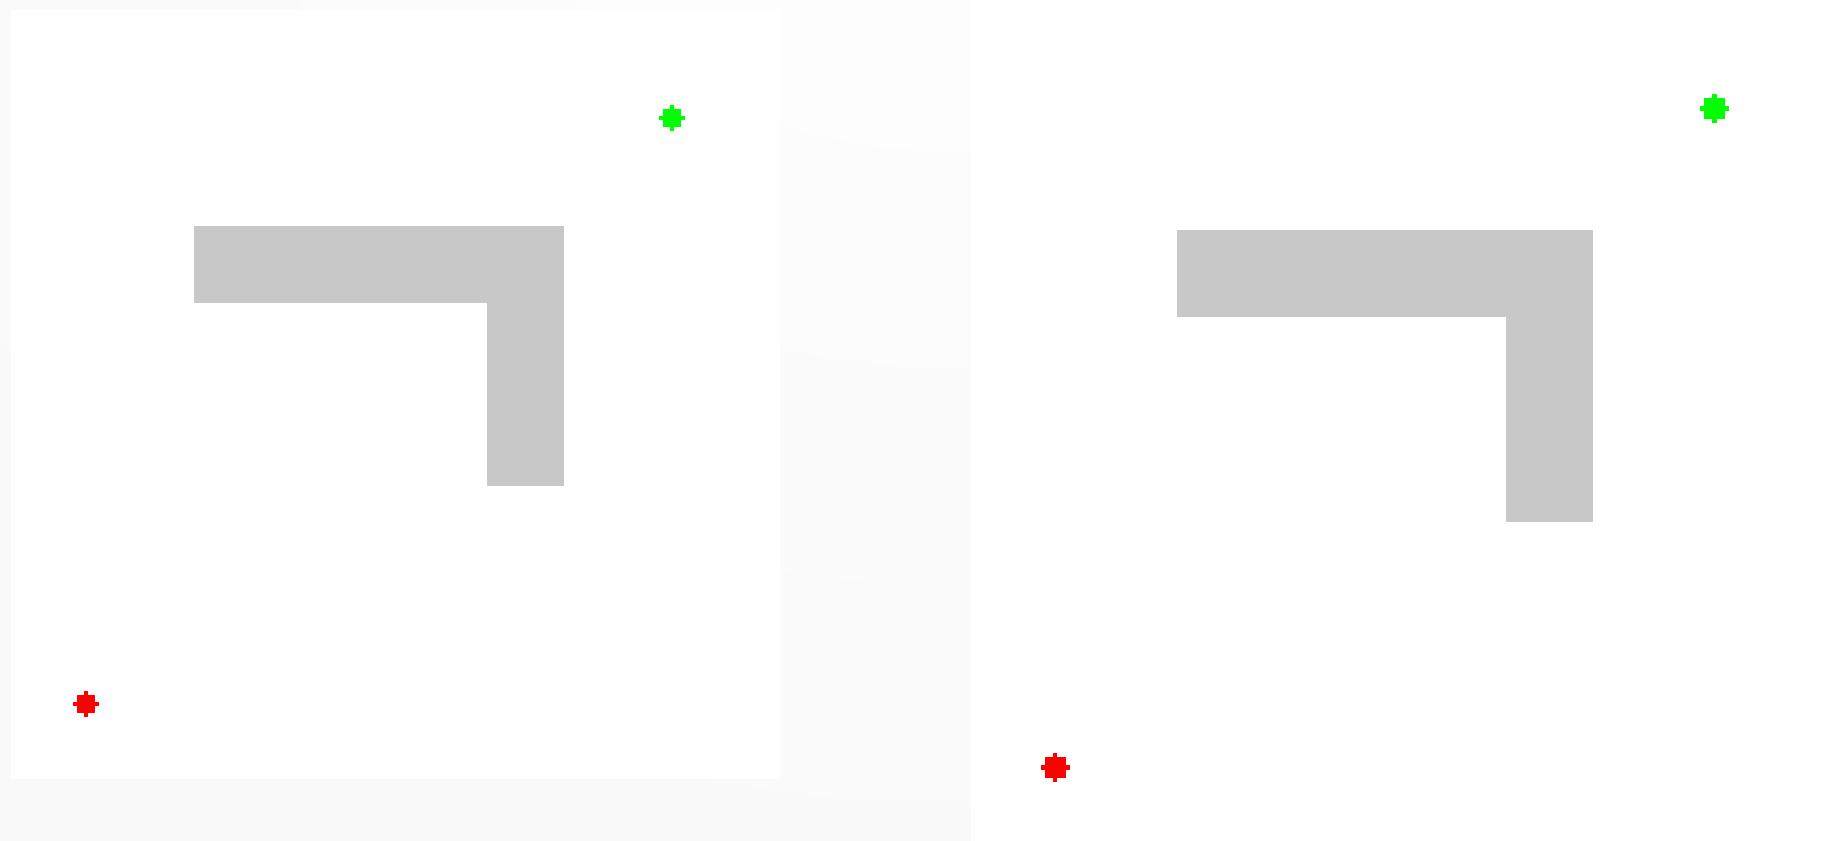
Dijkstra + heuristic: **A* algorithm** (*Hart, Nilsson & Raphael 1968*)



Preamble: path planning in a graph

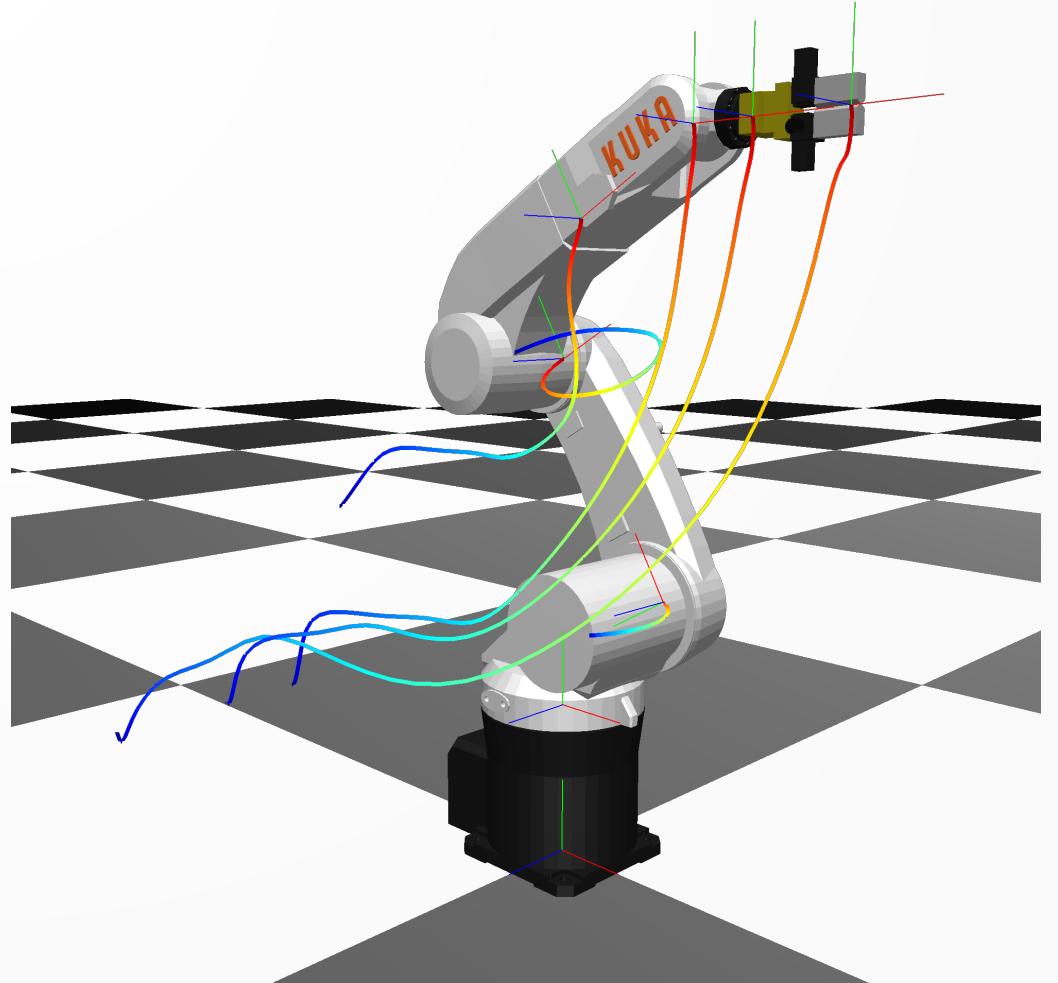
A*:

- A good heuristic is not always easy to find.
- Only efficient if the solution is not “too” complex.



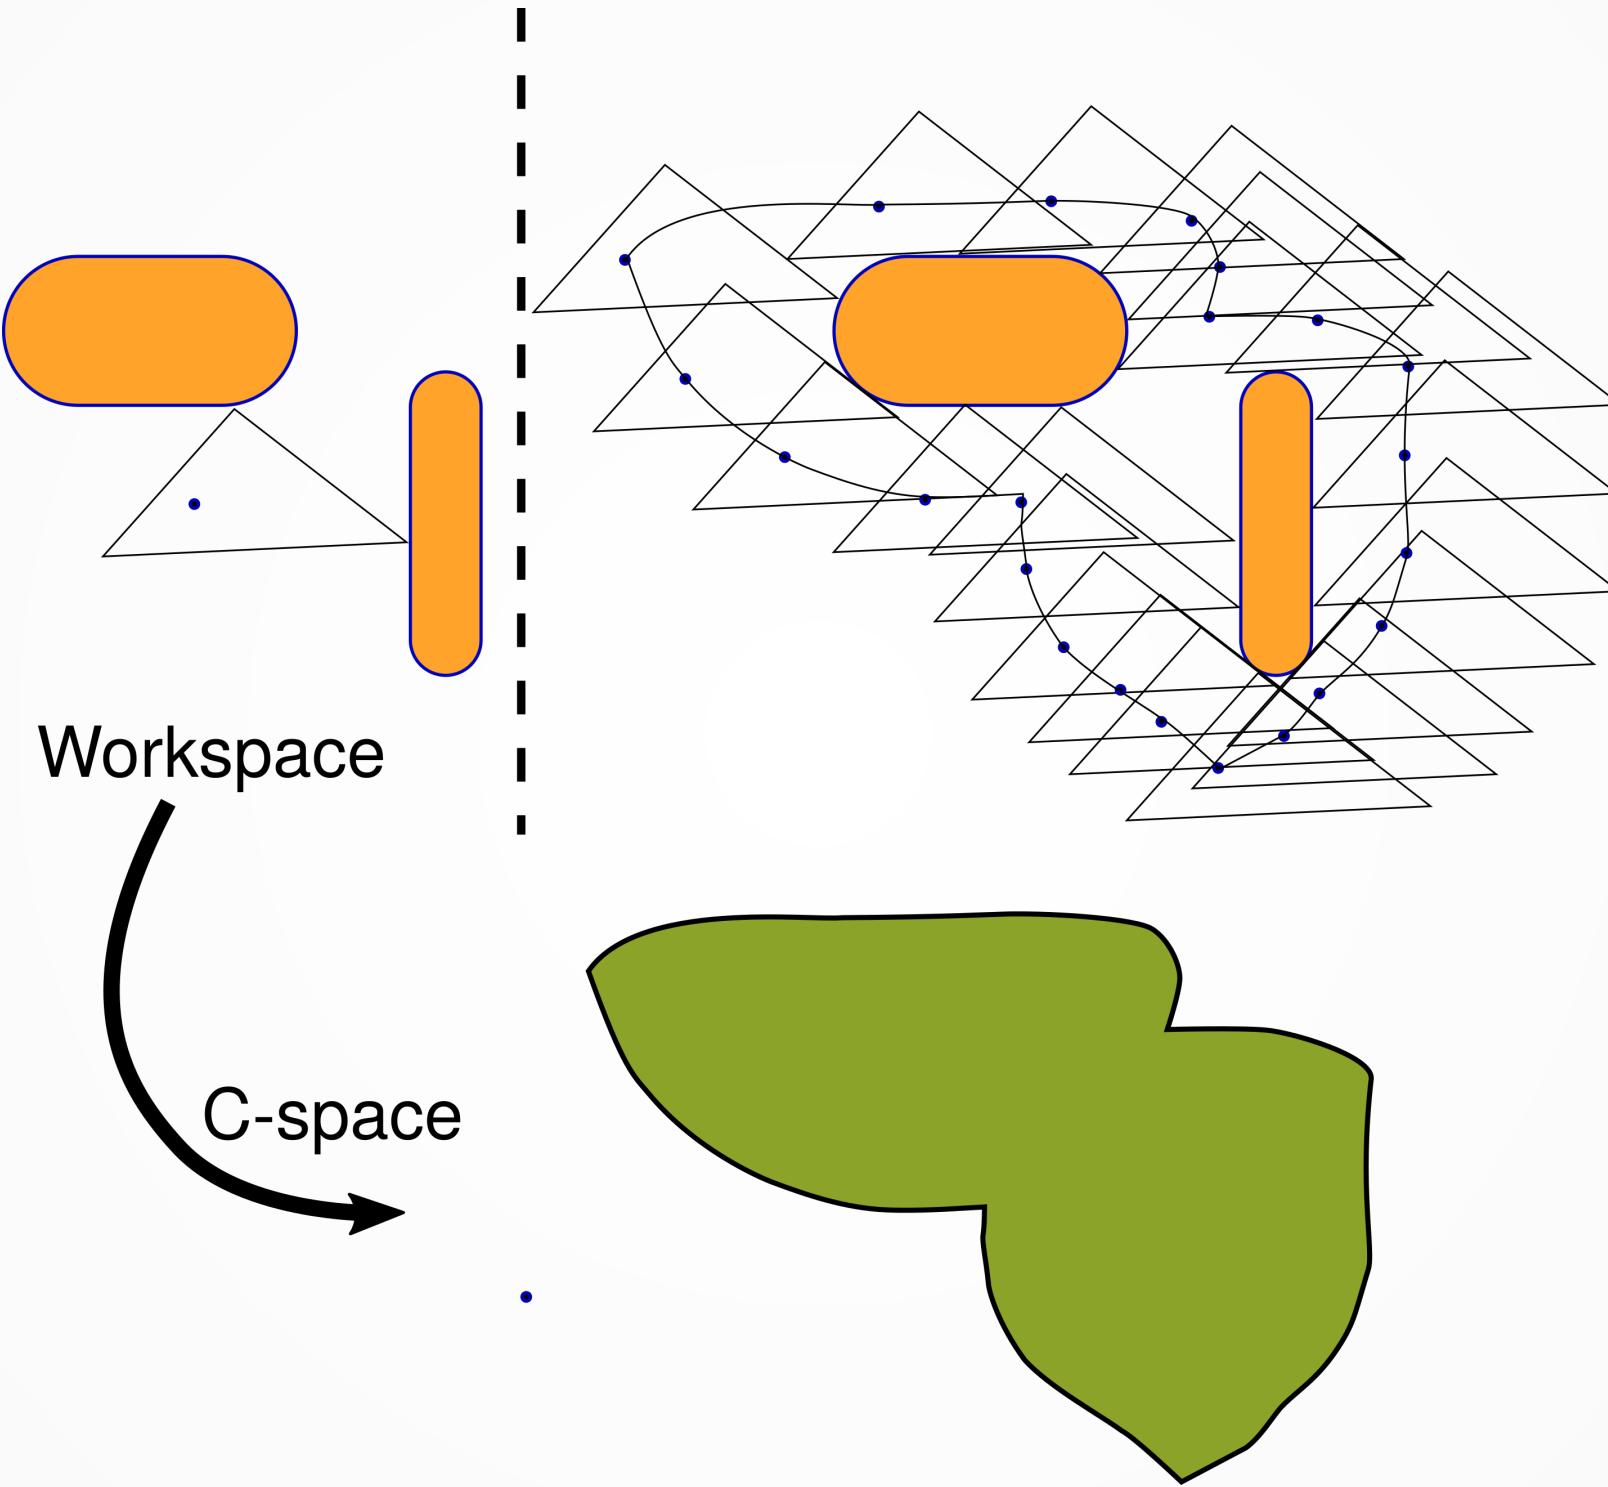
Challenges

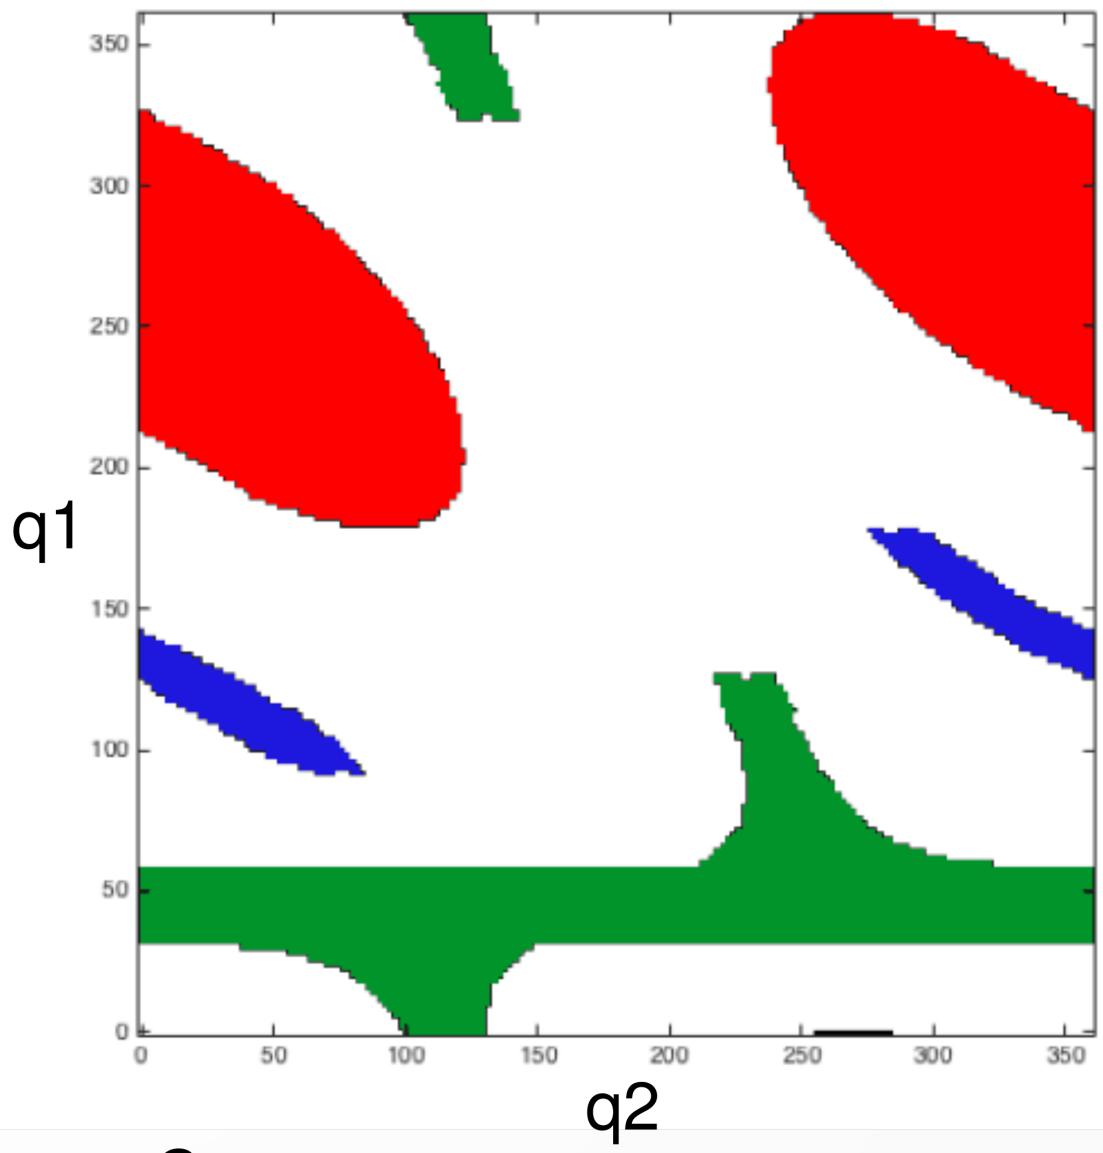
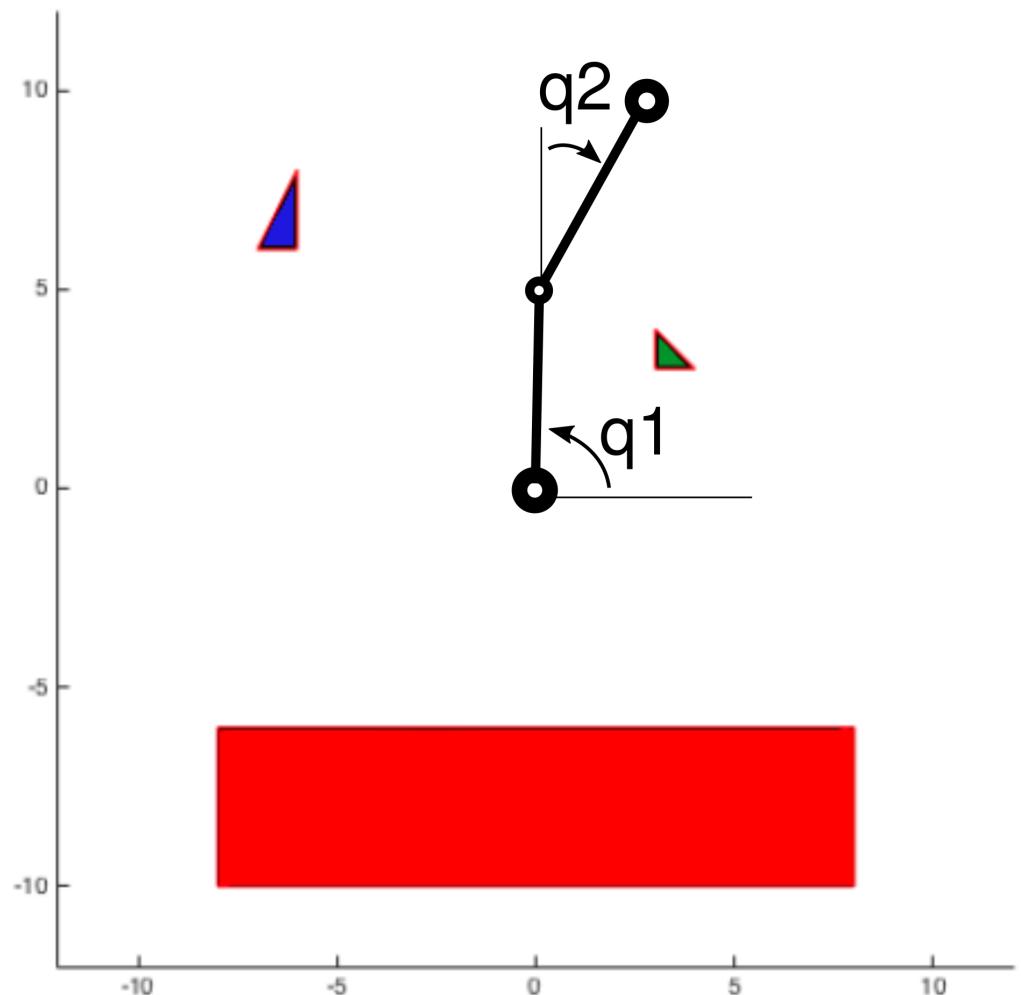
- Continuous,
high-dimensional
configuration spaces
- Continuous,
high-dimensional
action spaces



Challenges

- The configuration space, or C-space, is a space where, conveniently, the robot is just a point (an element of a vector space). Remark: we, humans, never think about the C-space.
- Obstacles are known in the **workspace**, not directly in the **configuration space**.



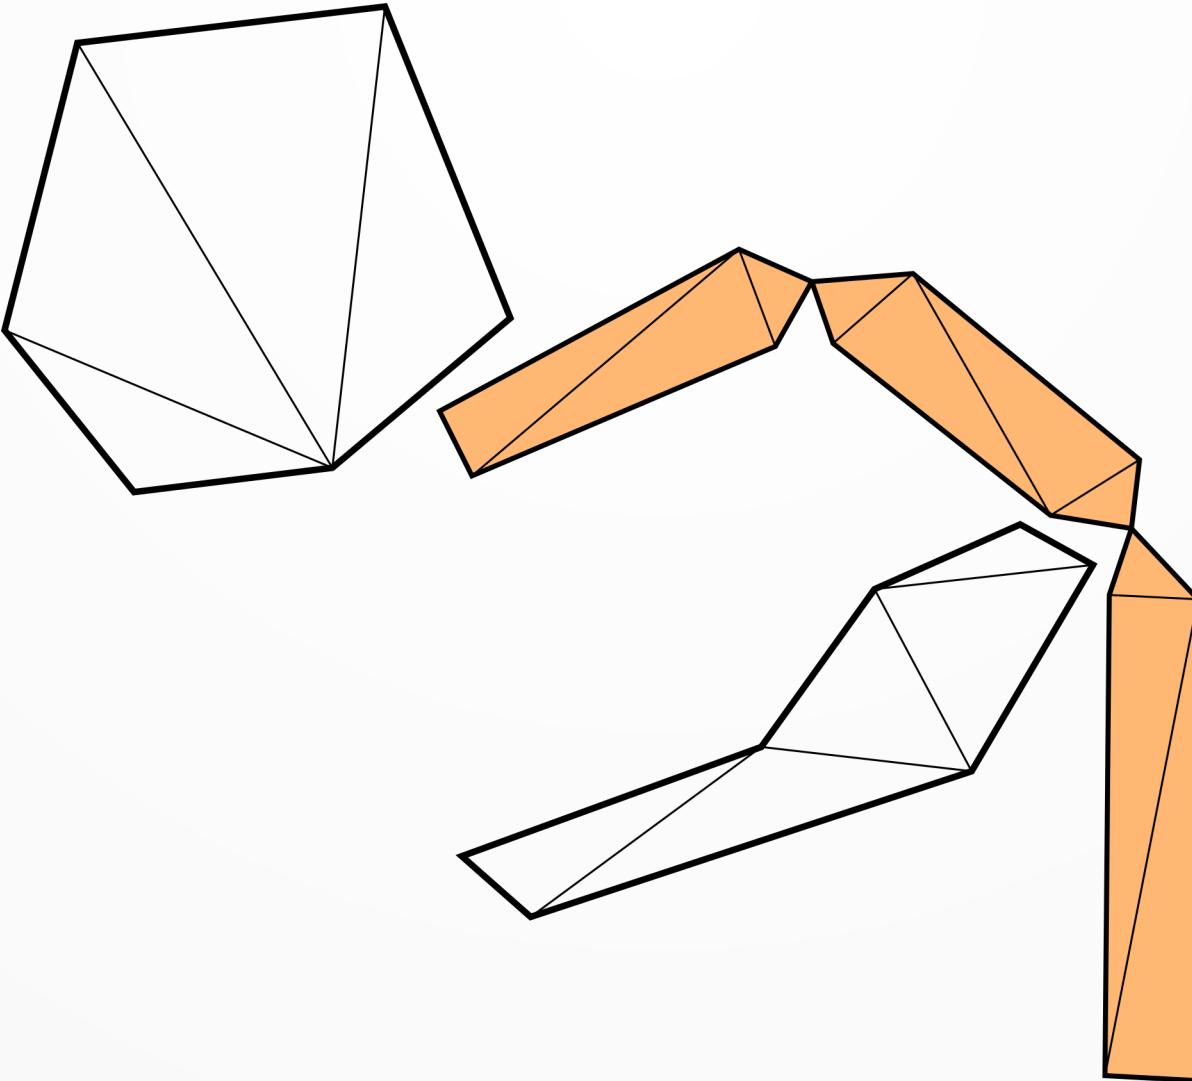


Workspace → C-space

Challenges

- No efficient way to know the obstacles in the C-space.
 - Most approaches perform motion planning “blindly” (without starting from the obstacles):
 - Configuration are first tried and then collisions are checked.
 - This leads to **“sampling-based”** motion planning, which relies on two main components:
 - (1) **efficient random exploration**
 - (2) **efficient collision checking**

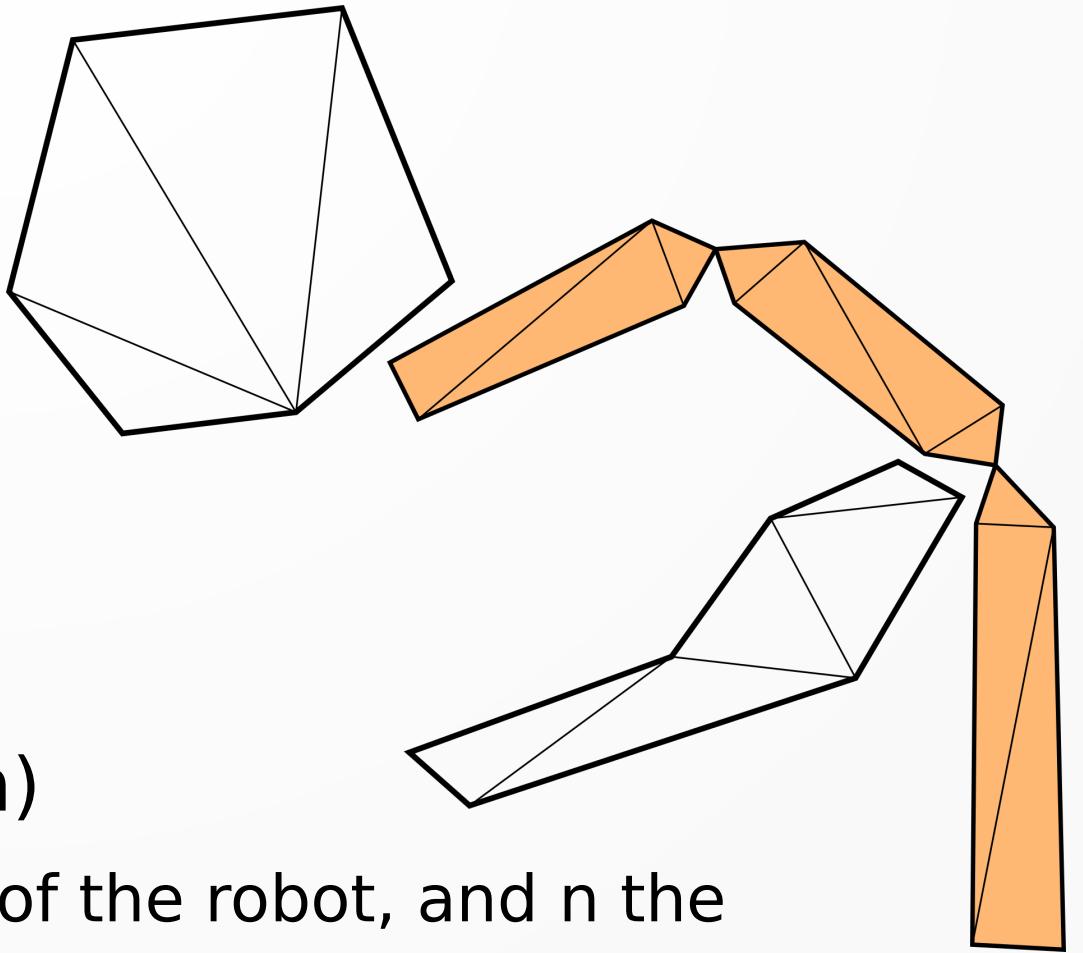
Collision checking



Triangulation: see for instance Delaunay triangulation (1924)

Collision checking

- Naive approach:
 - Test intersection for every possible pair of triangles
(that belong to different rigid bodies, one of these rigid bodies being part of the robot)
- Complexity: close to $n \times (n+m)$
(with n the number of triangles of the robot, and m the number of triangles in the environment)



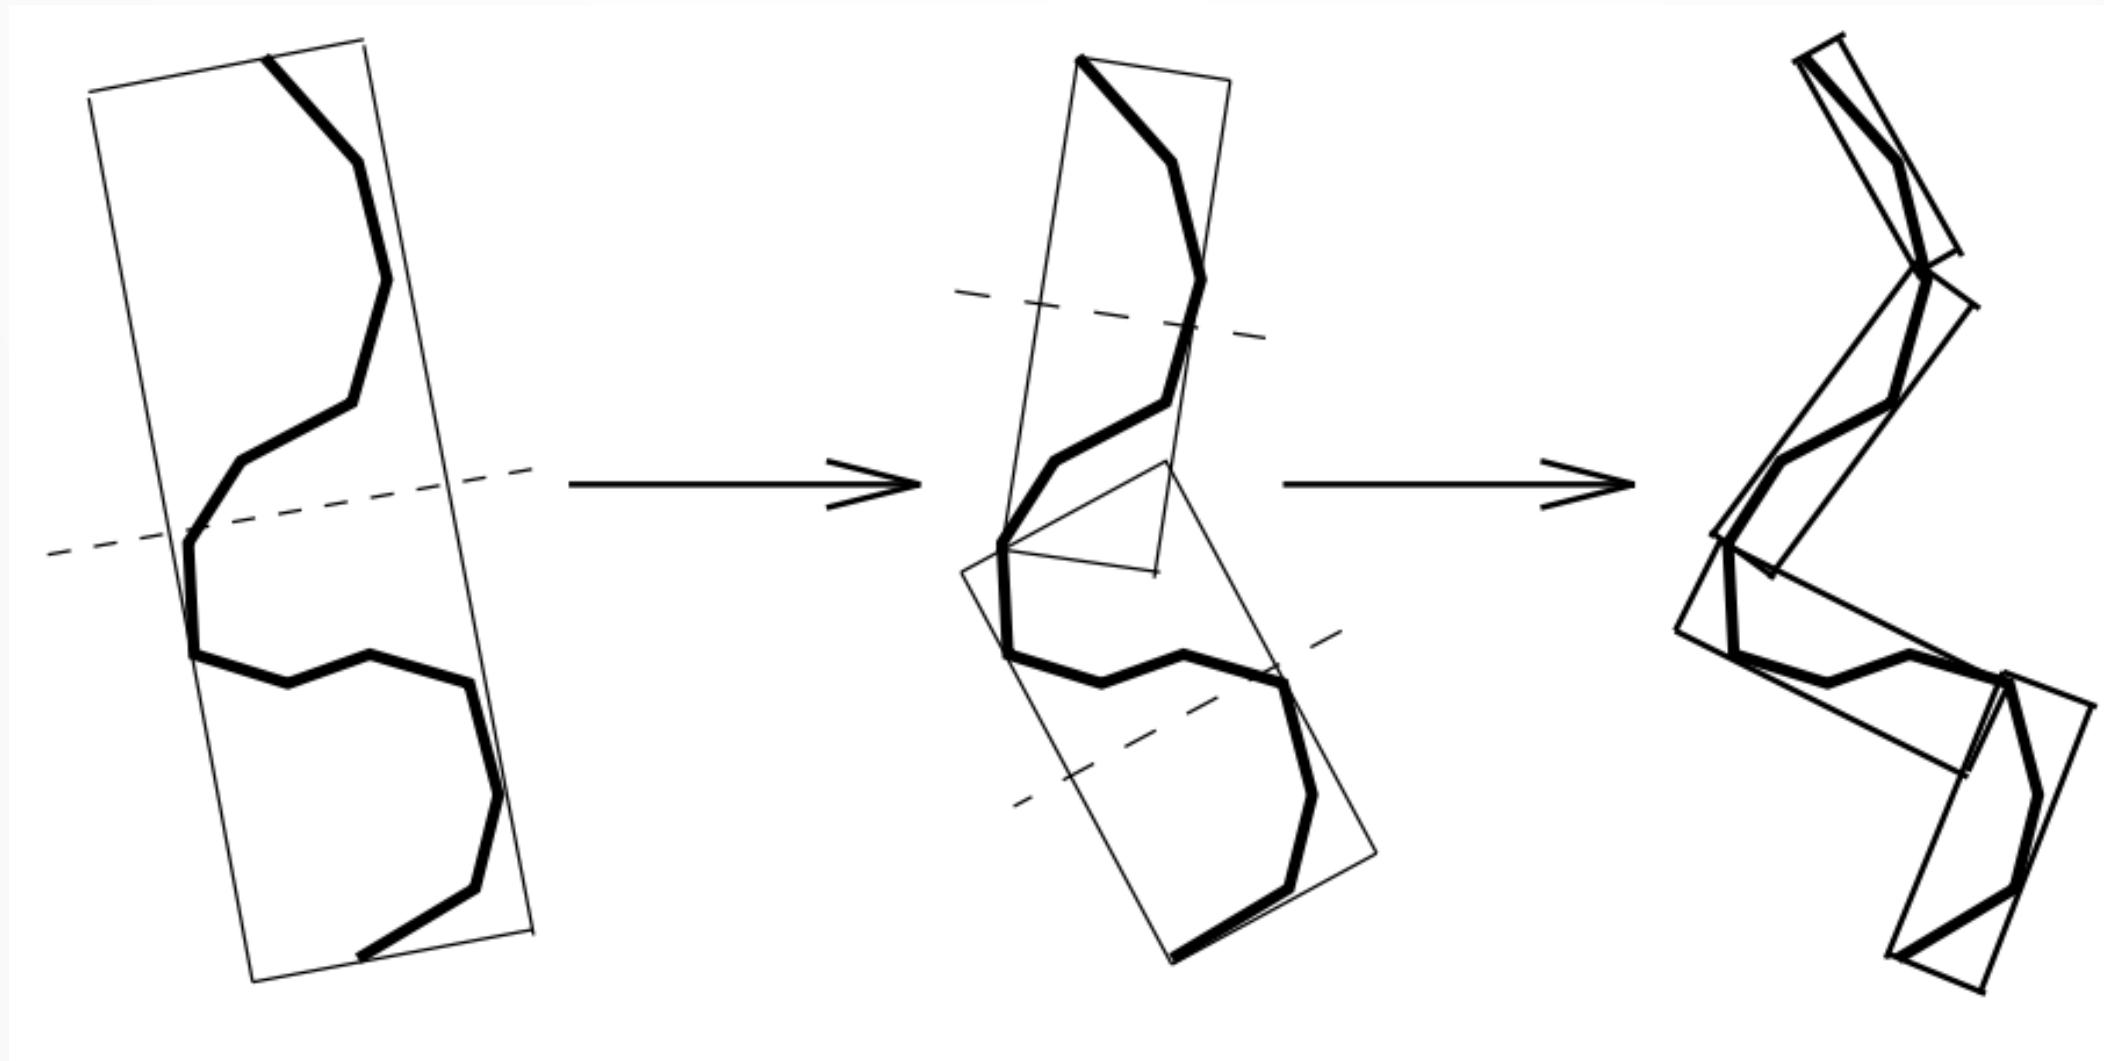
Collision checking

- Better approaches:
 - Incremental methods

E.g.: **A fast algorithm for incremental distance calculation**, M.C. Lin & J.F. Canny, IEEE ICRA 1990
 - Hierarchical methods

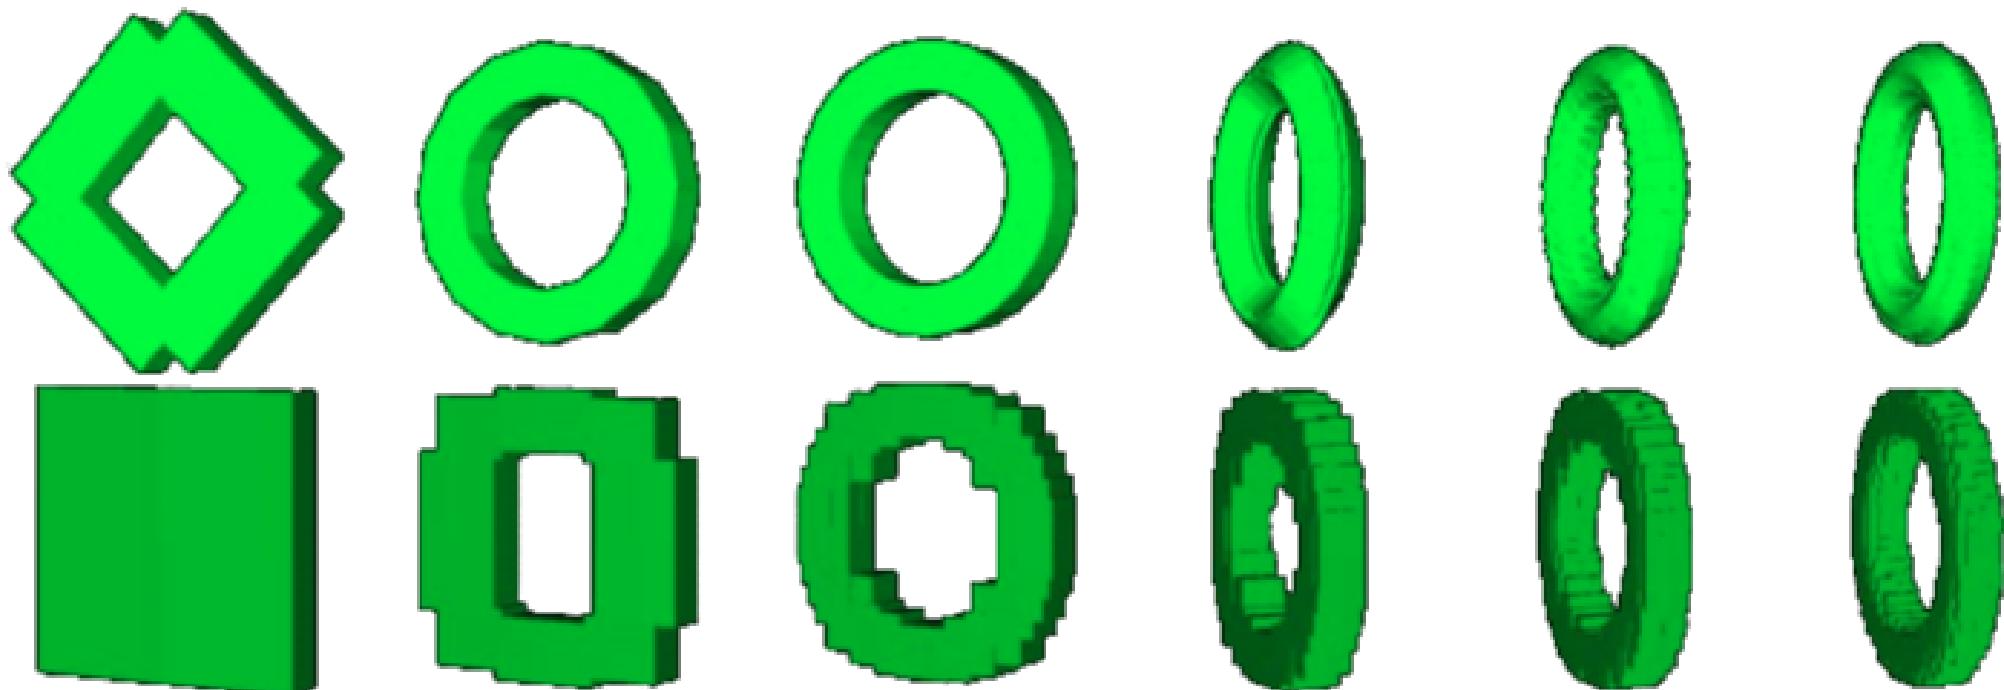
E.g.: **OBBTree: A Hierarchical Structure for Rapid Inference Detection**, S. Gottschalk, M.C. Lin, D. Manocha, ACM SIGGRAPH 1996

OBBTree



OBBTree

OBB vs. AABB (Axis-Aligned Bounding Boxes):

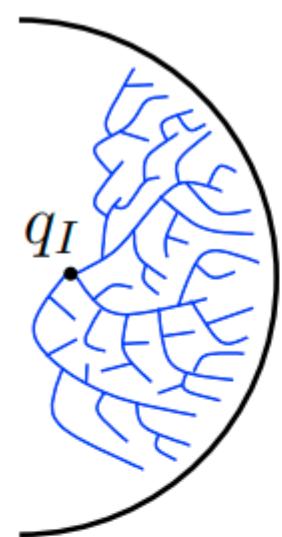


Sampling-based exploration

- To simply find configurations, one possibility is to use *low-discrepancy sequences*, which cover the search space quickly and evenly (much better than uniformly distributed random samples).
- However, motion planning usually aims at finding paths, not only new configurations. **The standard approach are based on the iterative expansion of random trees.**

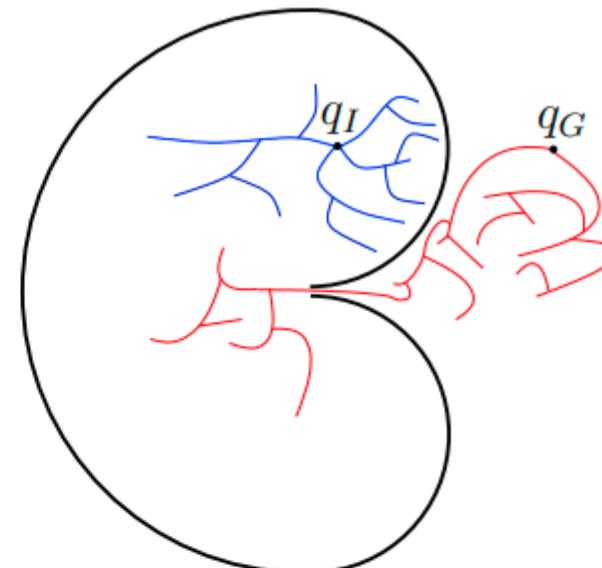
Single-query path planning

- Unidirectional methods: one random tree
- Bidirectional methods: 2 random trees
- Multi-directional methods: more than 2 trees

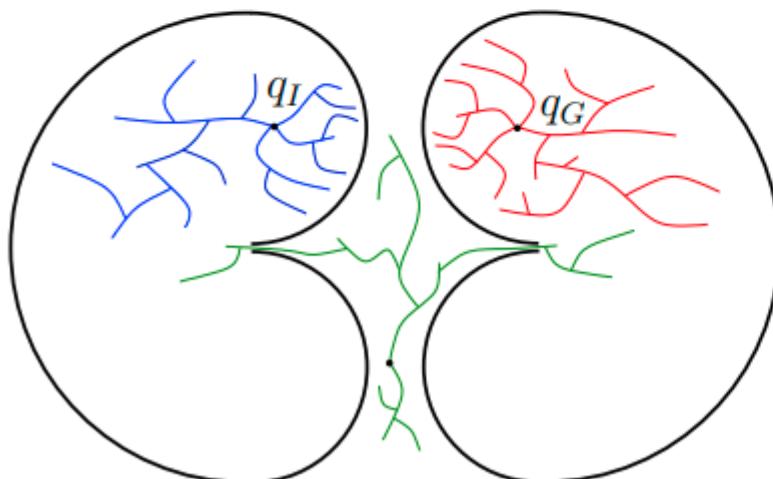


q_G

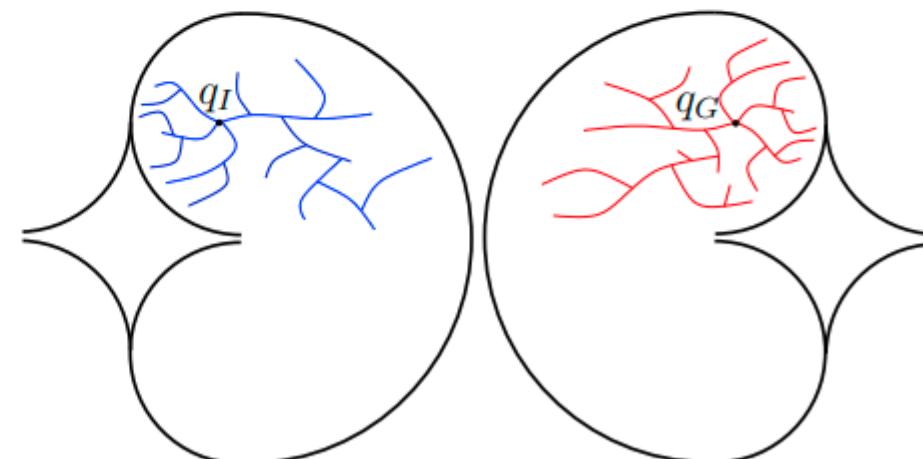
(a)



(b)



(c)



(d)

Selection-expansion

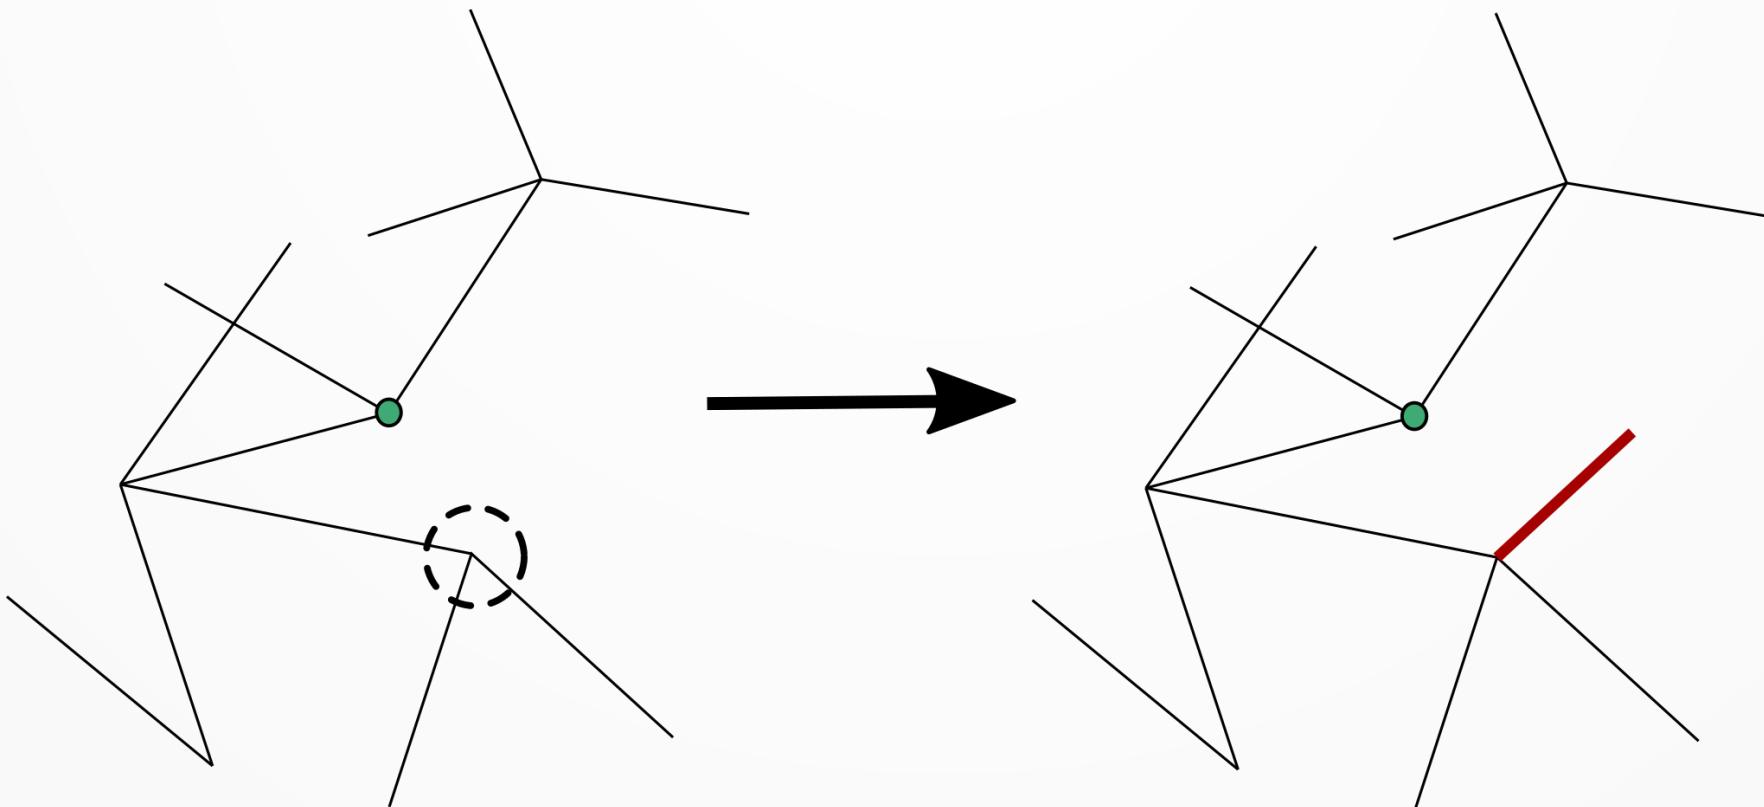
The structure of search algorithms based on trees is basically always the same, it consists in repeating the following operations:

- (1) **Selection:** select a node of the tree based on some criterion.
- (2) **Expansion:** try to “expand” from the selected node, i.e. construct a small collision-free motion from that node to a new node.

Remark: to check that a transition (edge) from a configuration to another is collision-free, we usually test regularly-spaced configurations along the transition (which requires a hand-tuned parameter: the delta between consecutive configurations).

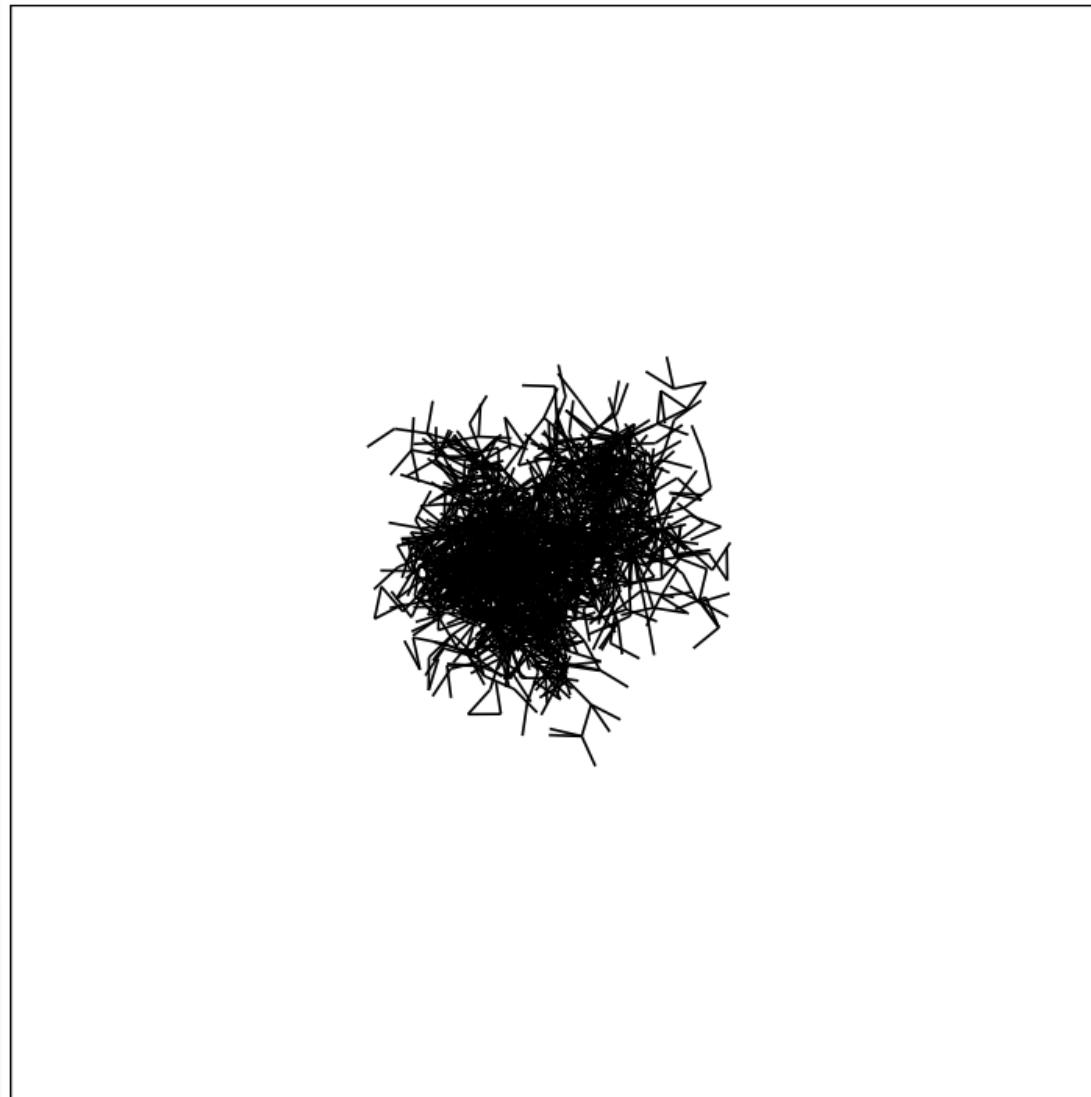
Naive random tree

Select a node randomly, and then expand it randomly.

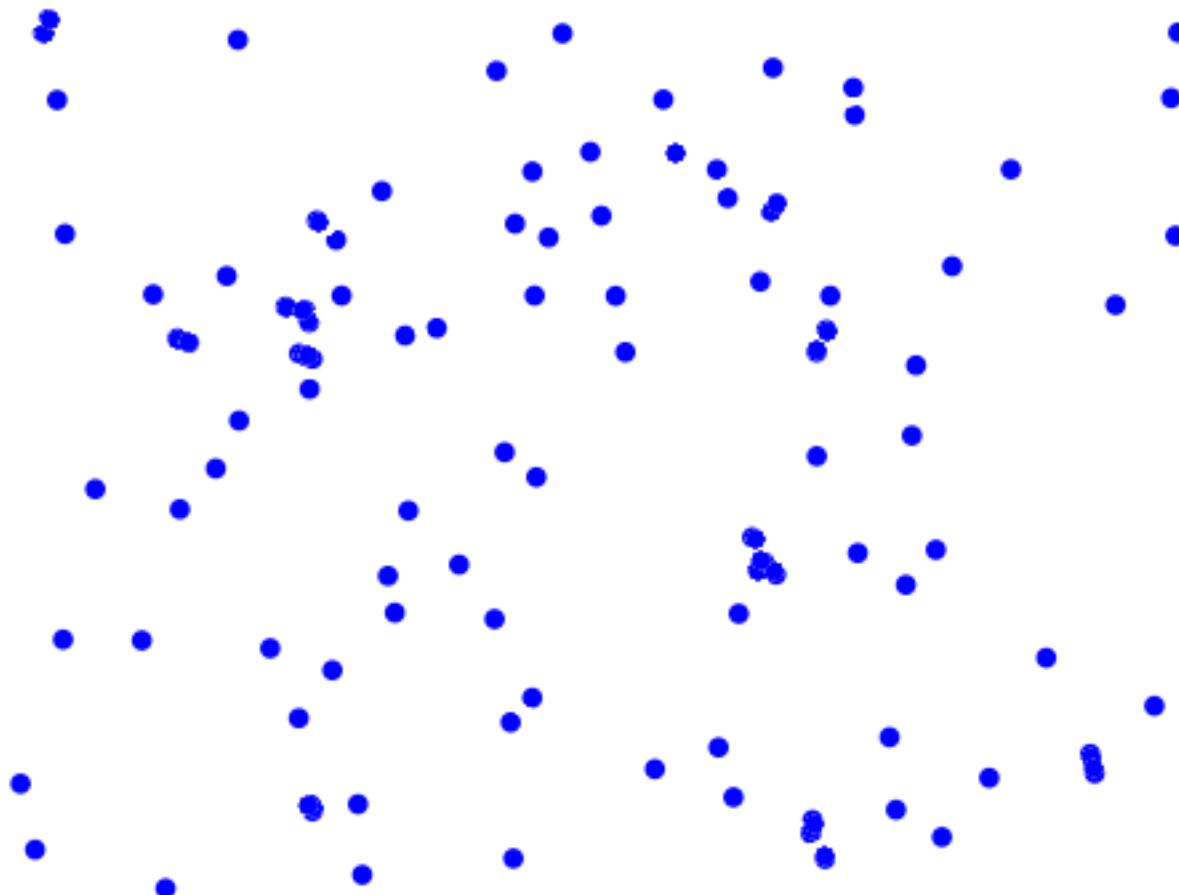


Naive random tree

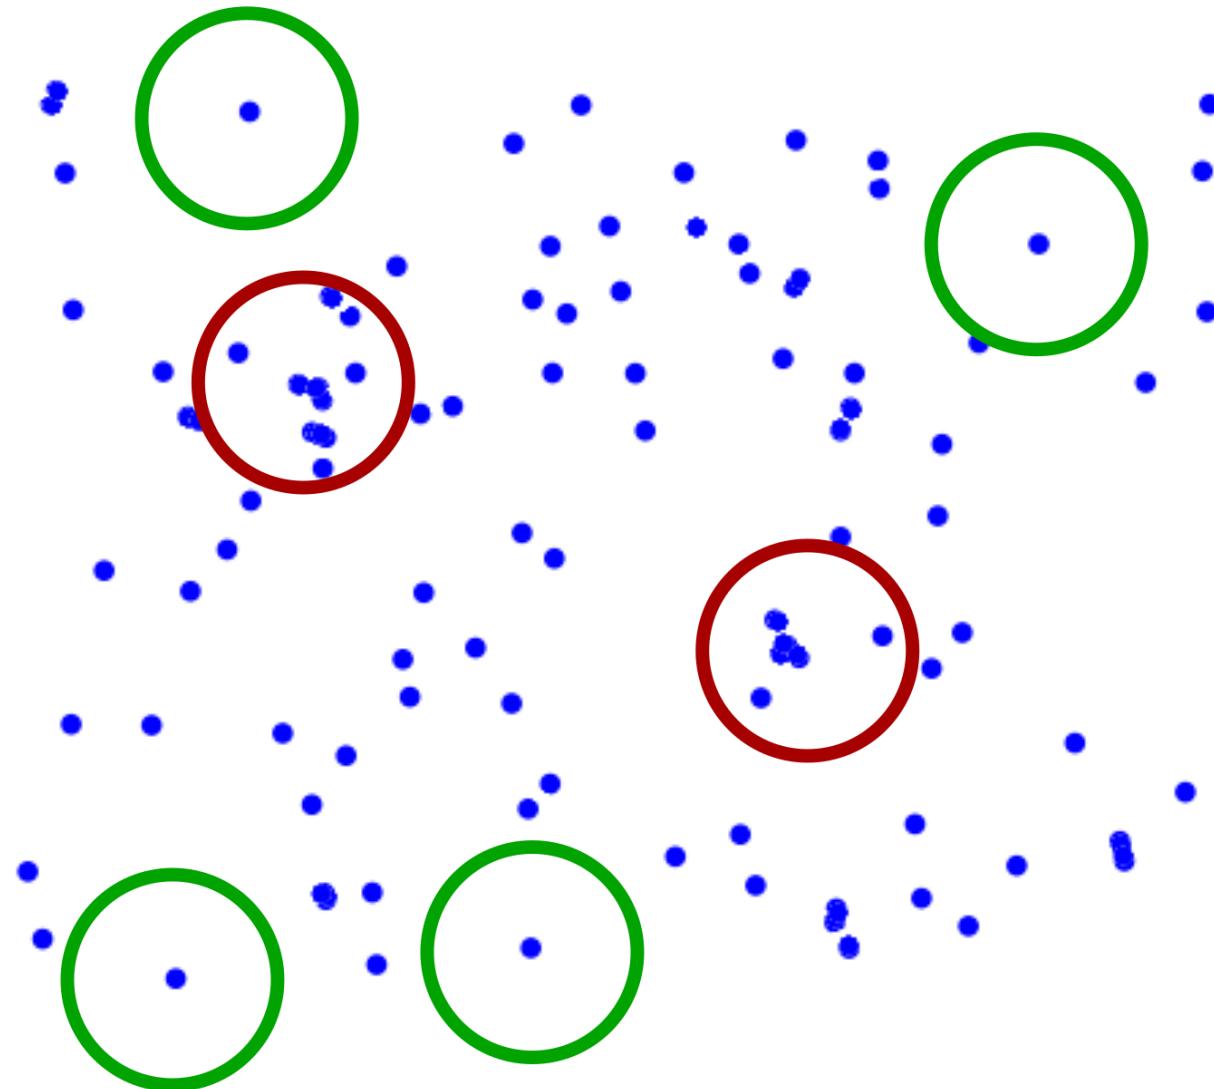
If we iterate, we end up with something like this (very dense):



Selection based on density



Selection based on density



Expansive Spaces Trees (EST)

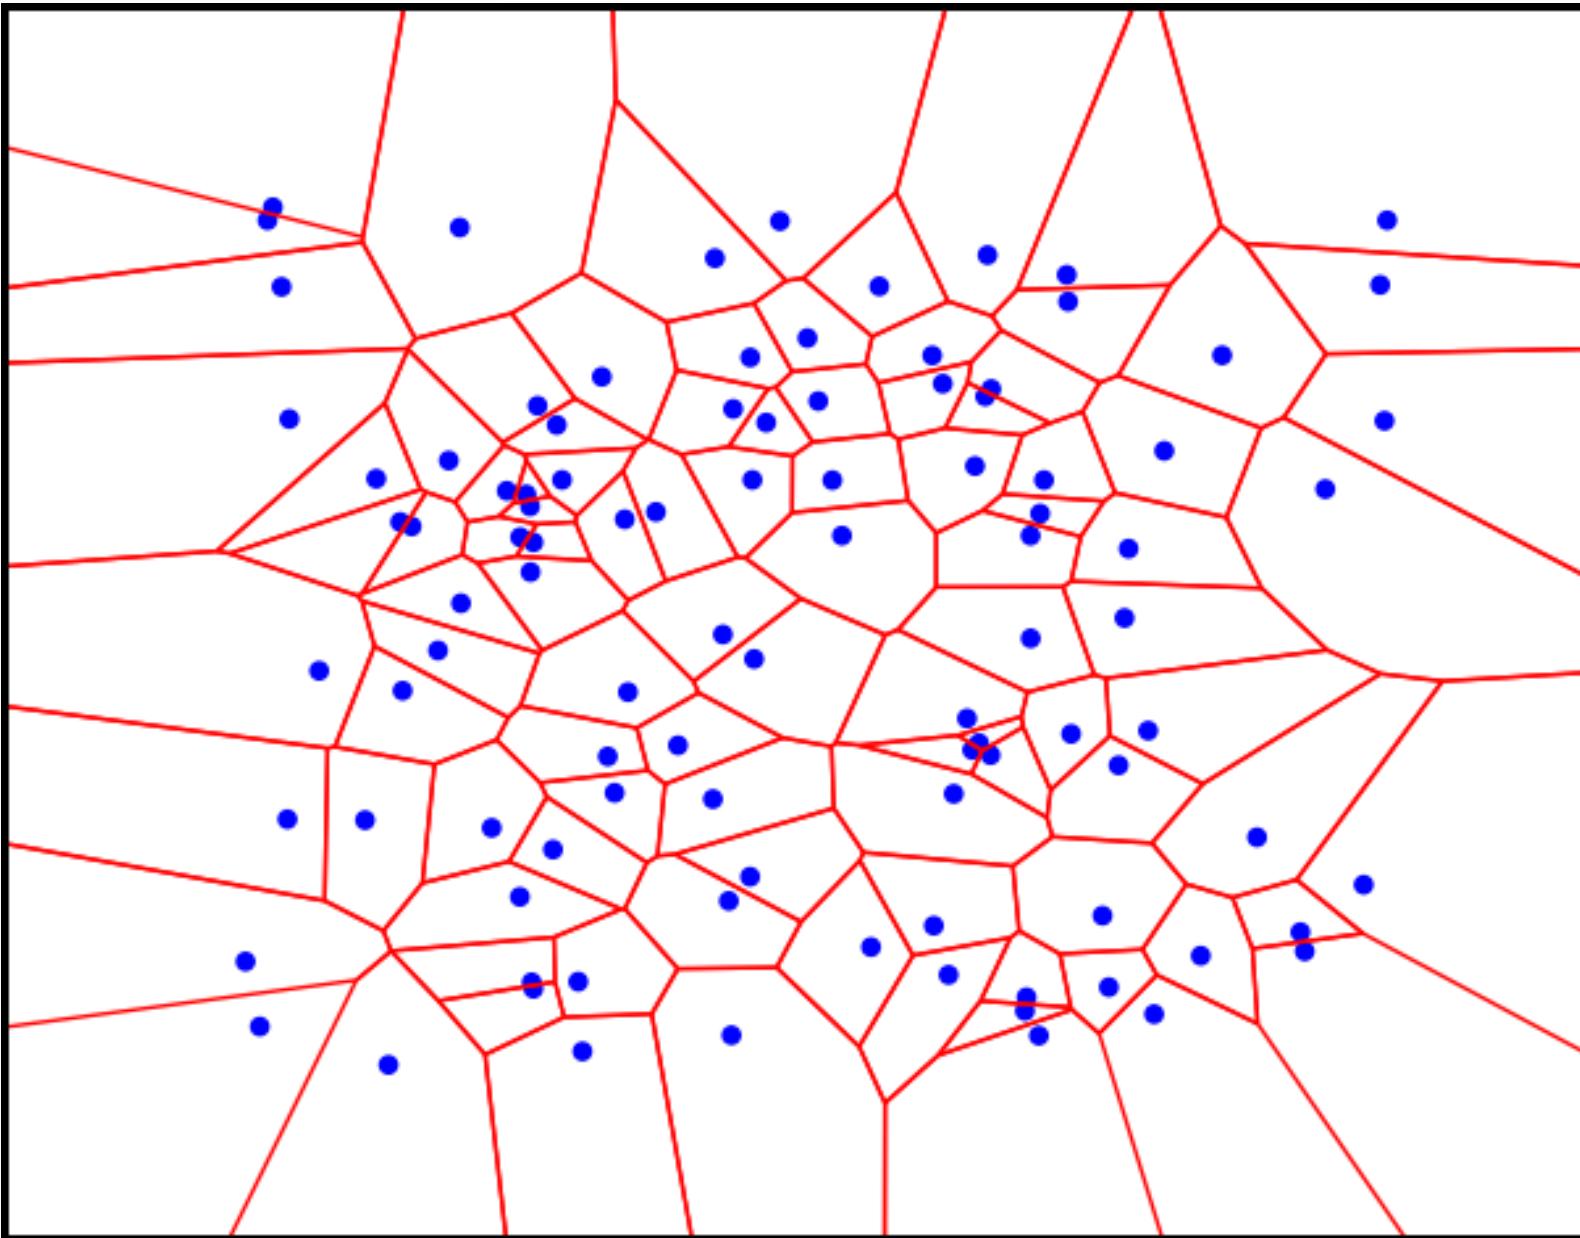
- When adding a new node, count its neighbors within an arbitrarily fixed distance to estimate the “local density” of nodes around it (= the number of nodes found divided by the volume). *Remark: update also the estimated density for the neighbors.*
- Nodes are selected with a probability that is inversely proportional to the estimated local density.

Ref.: **Path Planning in Expansive Configuration Spaces**, D. Hsu, J.-C. Latombe, R. Motwani, IEEE ICRA 1997.

Limitations of EST

- Requires good **nearest neighbors algorithm**. *The standard way to do this is to use k-d trees.* See **Multidimensional binary search trees used for associative searching**, J. L. Bentley, Communications of the ACM 1975.
- Covering with constant density fills holes, it does not attempt to expand as fast as possible.
- Once nodes are selected, expansions are too random.

Voronoi Diagrams

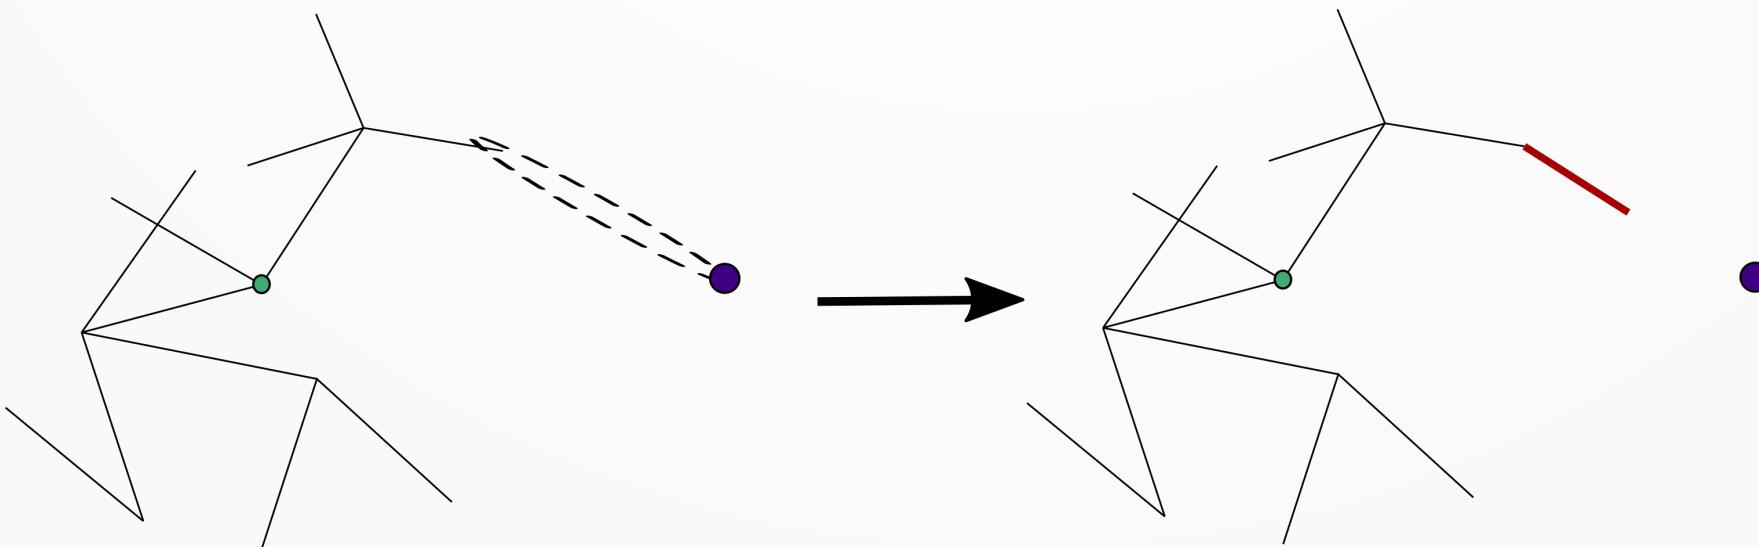


Voronoi Diagrams

- We would like to expand the node with the largest Voronoi cell, however this is very difficult to compute.
- *But, what if we draw with uniform probability a random sample in the configuration space, and then select the nearest neighbor within the nodes ? Doing so, nodes have a probability to be selected that is proportional to the volume of their Voronoi cell.* It results in a sampling of nodes with a bias towards large Voronoi cells.

Rapidly-exploring Random Trees (RRT)

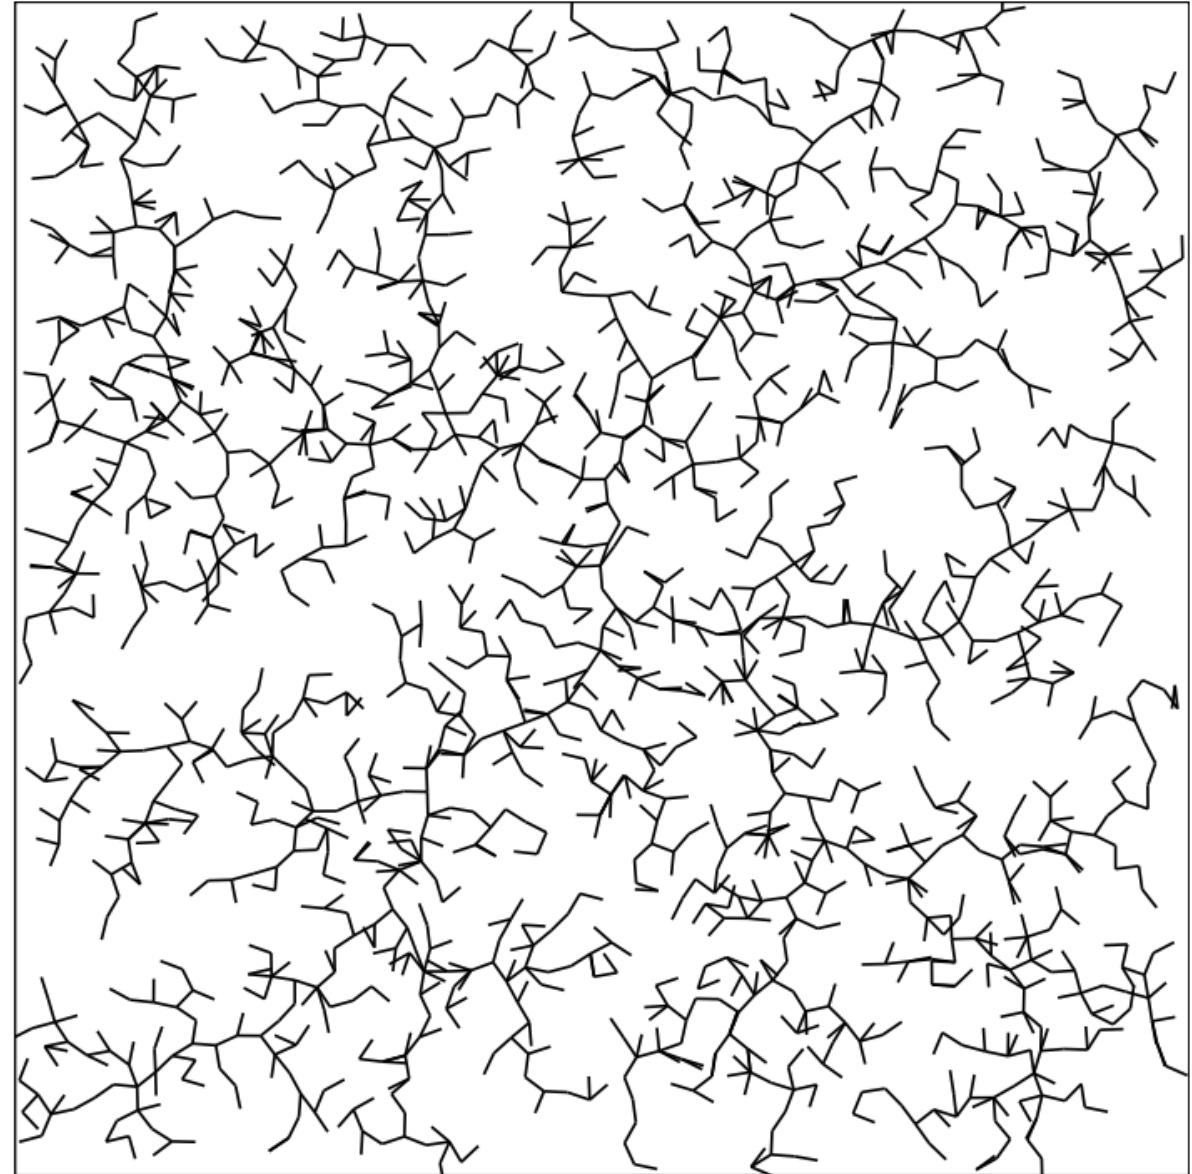
- Draw a random sample with uniform probability in the configuration space, and select the nearest neighbor within the nodes of the tree.
- Expand the selected node **toward** the randomly sampled configuration.



Rapidly-exploring Random Trees (RRT)

- See **Randomized Kinodynamic Planning**, S.M. LaValle & J.J. Kuffner, International Journal of Robotics Research, 2001.

Rapidly-exploring Random Trees (RRT)



Limitations of RRT

- Also requires a lot of nearest neighbor computations.
- Need to know the bounds of the configuration space →
In fact, there are simple solutions to that (example: use
*bounds ~20% larger than the smallest bounding box of
the current nodes*).
- **What if sampling a random configuration is hard?**
- **What if expanding toward a sample is hard?**
- **It is asymptotically complete, but not
asymptotically optimal.**

Limitations of RRT

- If there exists a path from A to B (B a region of strictly positive volume), then RRT will find it with probability one (*if given enough time*).
- However, it will not find an *optimal path* (for example the straight line if is collision-free).

RRT*

RRT* is a variant of RRT that includes some “rewiring” of the tree to achieve asymptotic optimality.

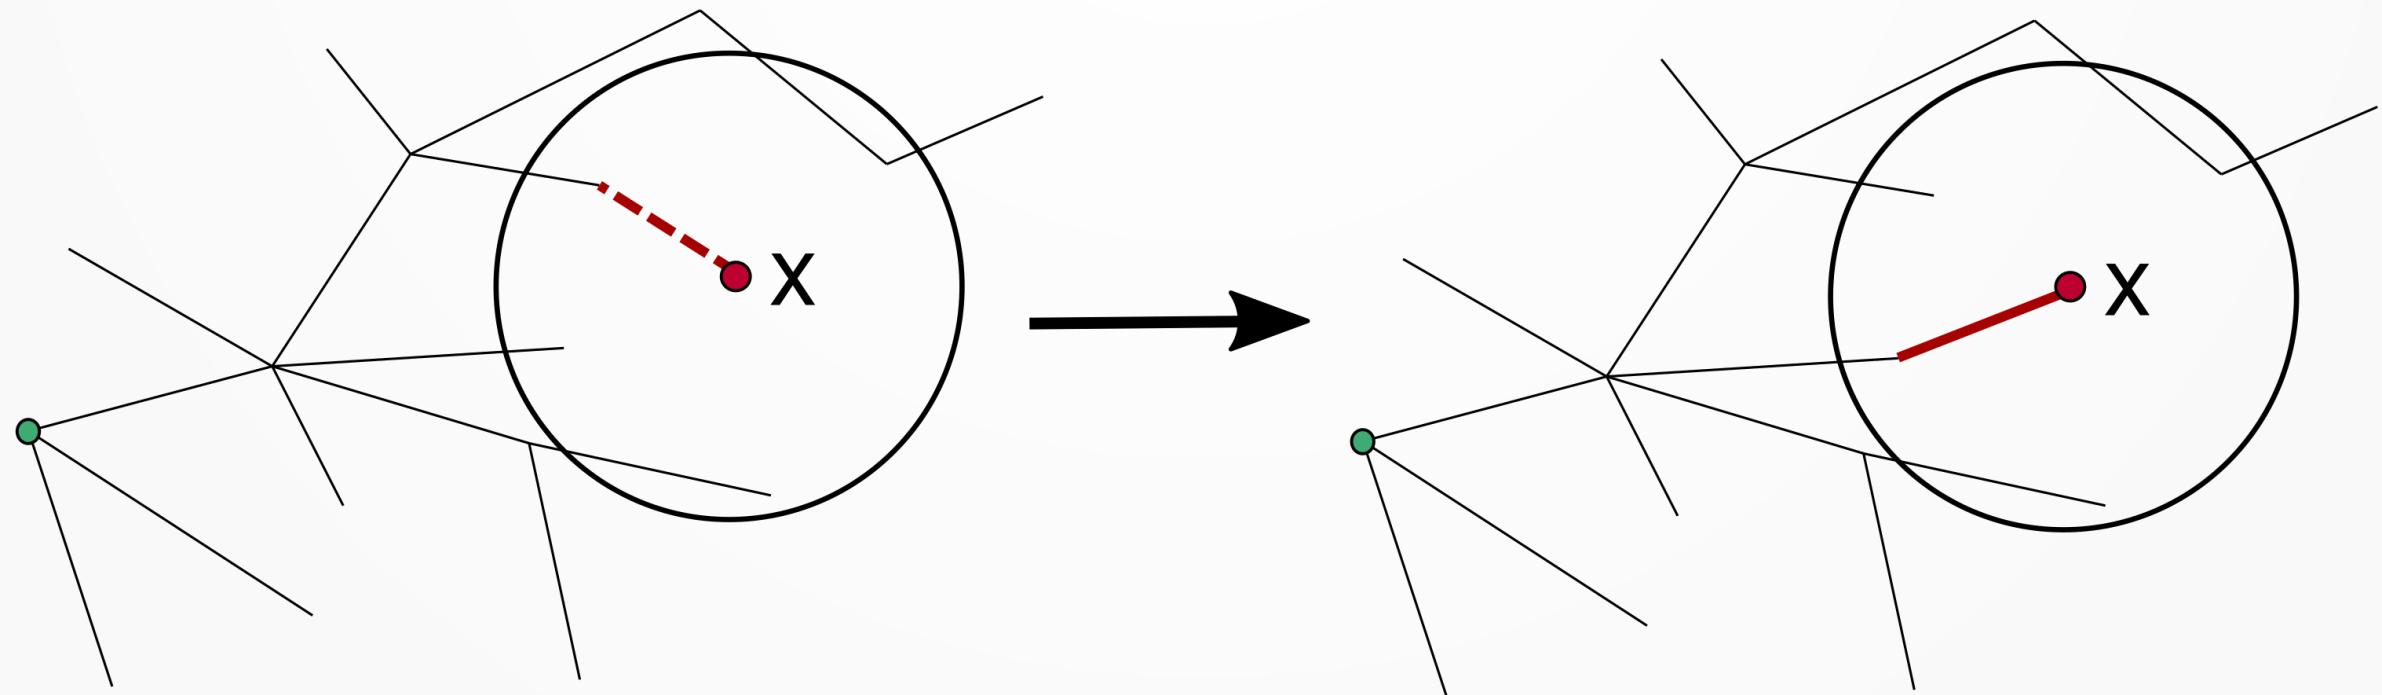
See: **Sampling-based algorithms for optimal motion planning**, S. Karaman & E. Frazzoli, International Journal of Robotics Research, 2011.

RRT*: the algorithm

(1)

- Run an RRT selection-expansion, it adds a new node x to the tree. Consider the neighbors of x in the tree within a sphere of radius r (*remark: r depends on the current total number of vertices in the tree*)
- Modify the expansion by choosing a parent for x that minimizes the cost of the total path from the root to x .

RRT*: the algorithm

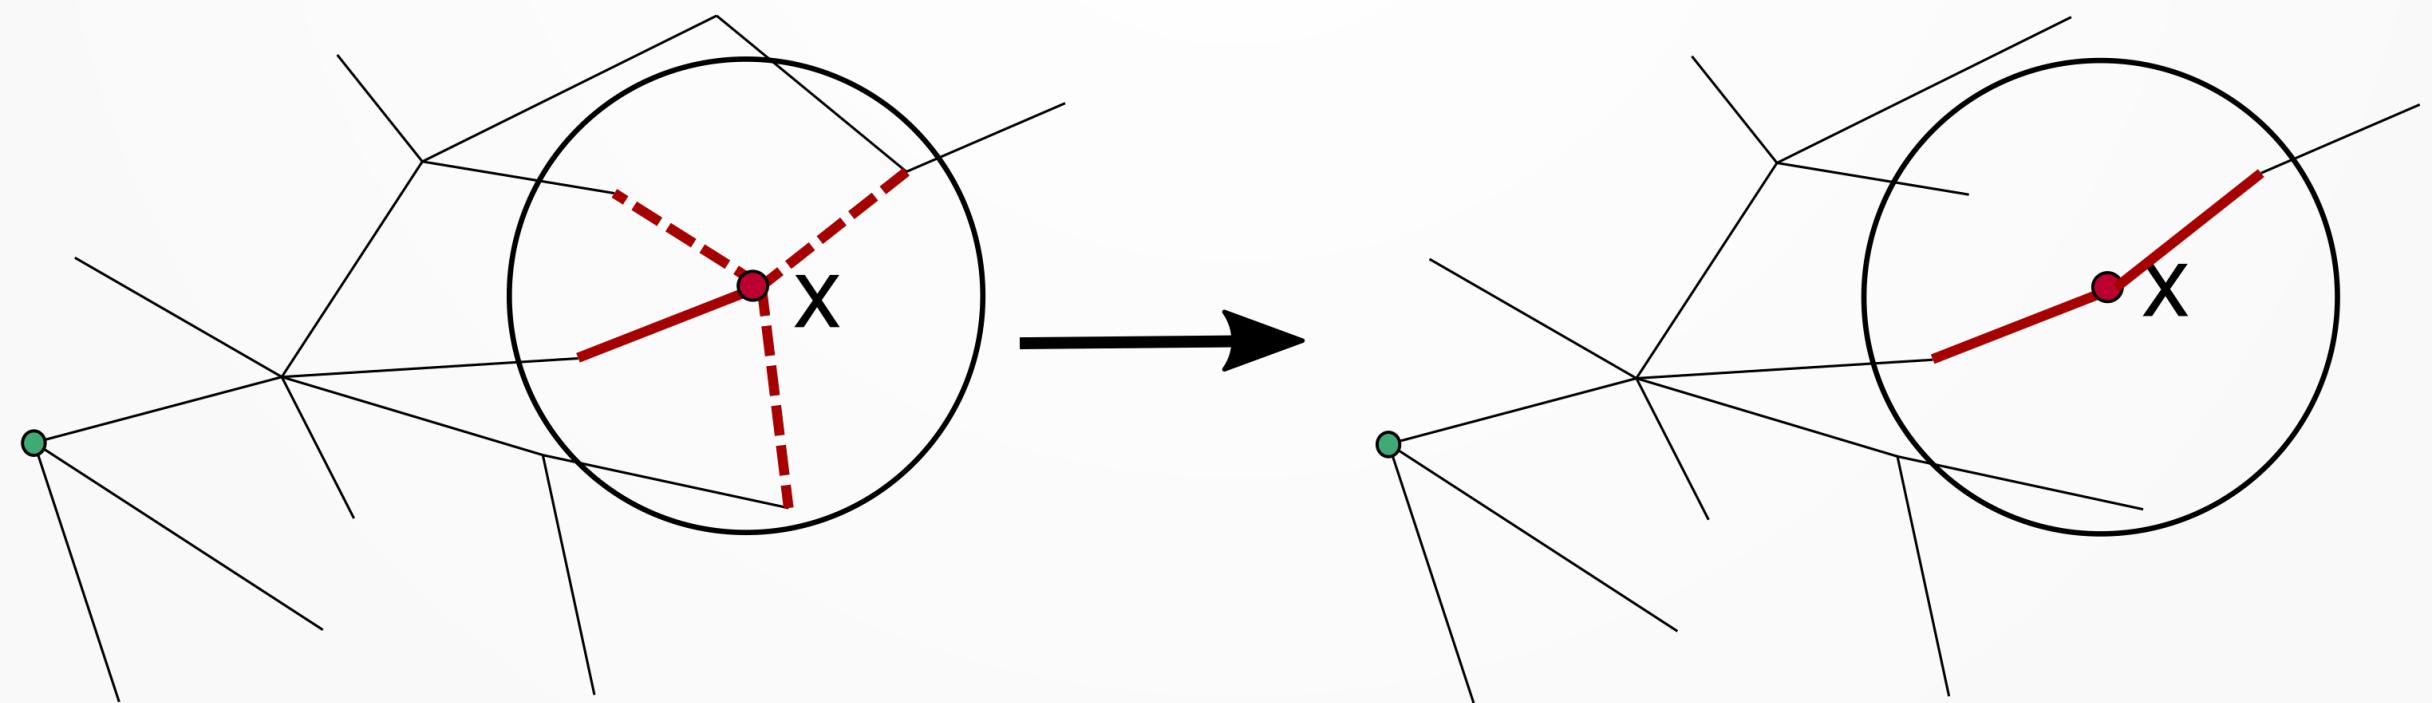


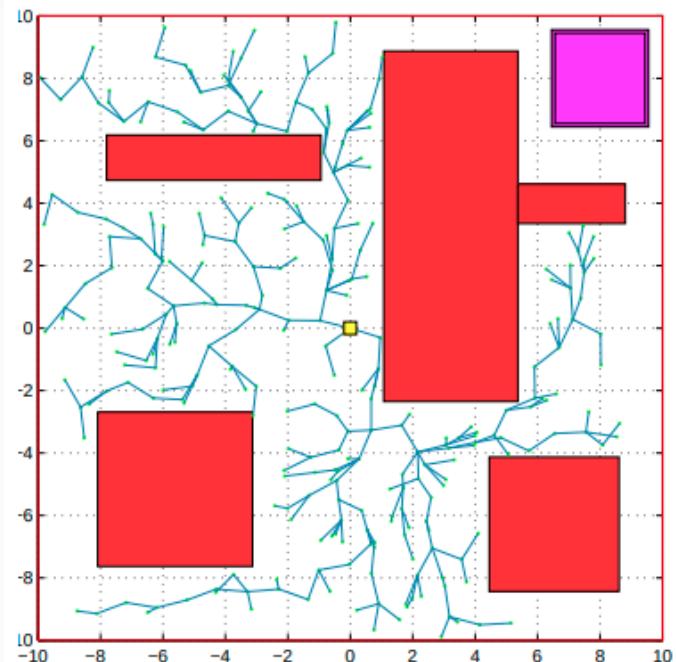
RRT*: the algorithm

(2)

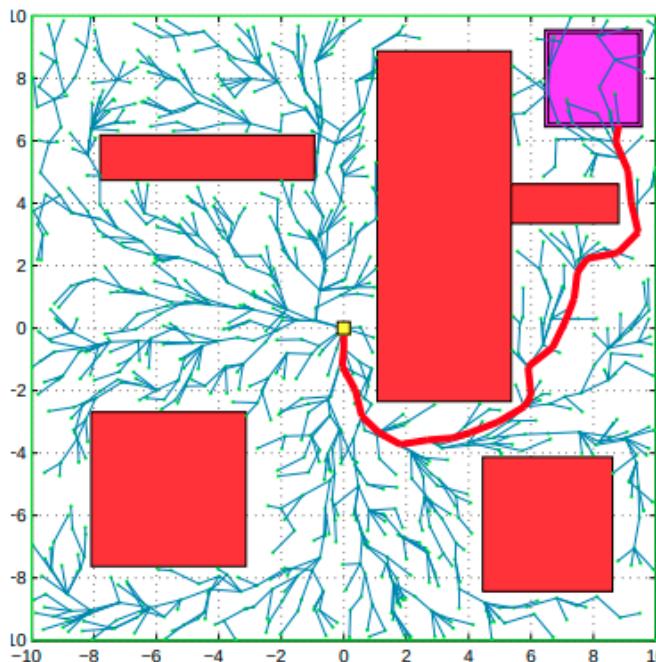
- Consider edges from x to all the “neighbors” y in the sphere, and add an edge to a neighbor y if it creates a path to y that is shorter than the currently existing path. If an edge to y is added, the edge to y from its previous parent is removed.

RRT*: the algorithm

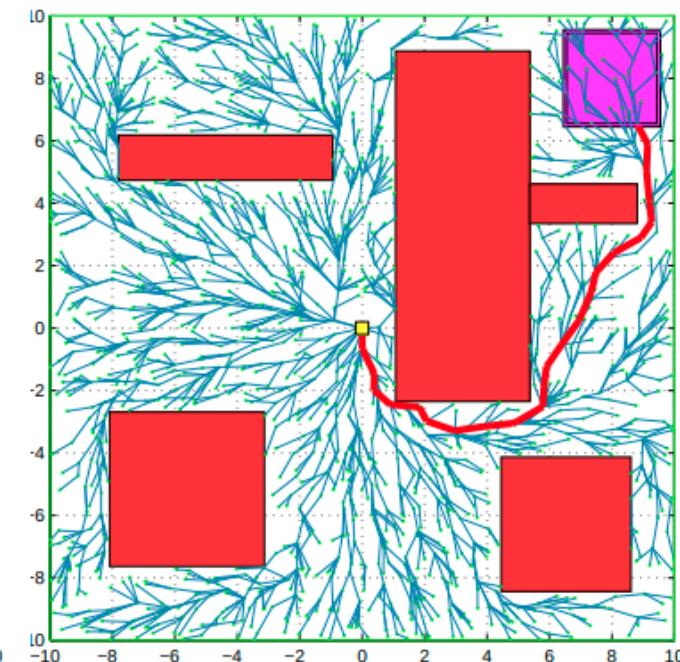




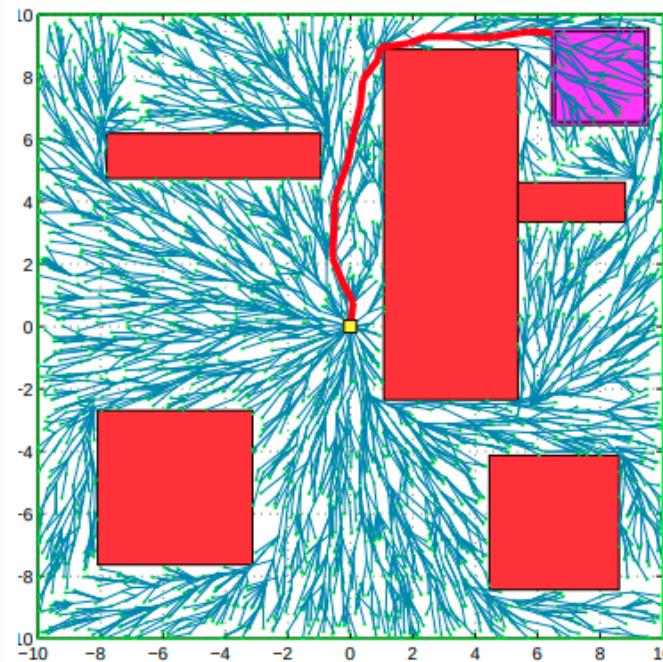
(a)



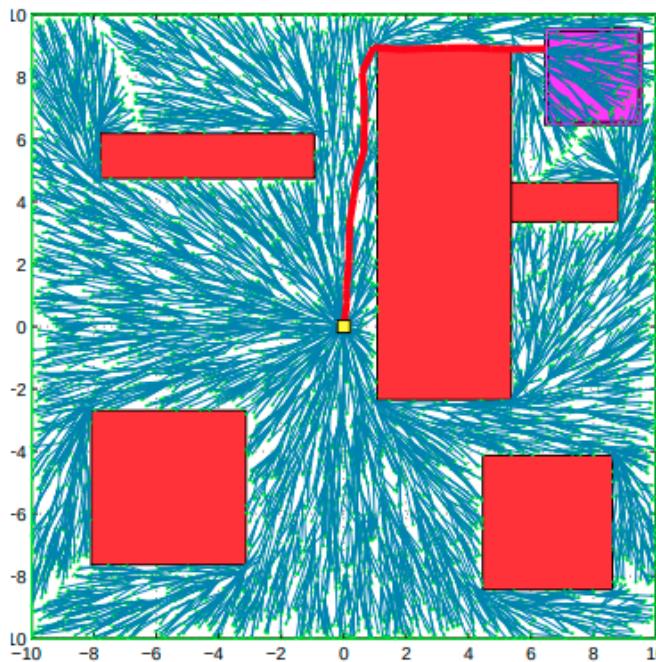
(b)



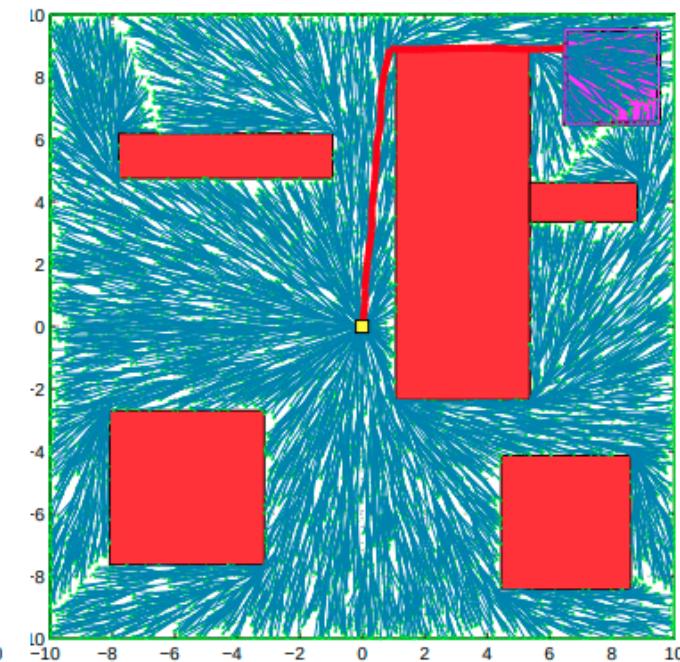
(c)



(d)



(e)



(f)

Non-holonomic and kinodynamic motion planning

In motion planning problems with non-holonomic or kinodynamic constraints, not all geometric trajectories are feasible. In such cases, the planning must be done in the state space, not only in the configuration space (remark: typically, the state includes velocities).

In general, via discretization, the reachable states \mathbf{x}' from a state \mathbf{x} verify :

$\mathbf{x}' = \mathbf{x} + \mathbf{f}(\mathbf{x}, \mathbf{u})$, for some possibly constrained control input \mathbf{u} .

Non-holonomic and kinodynamic motion planning

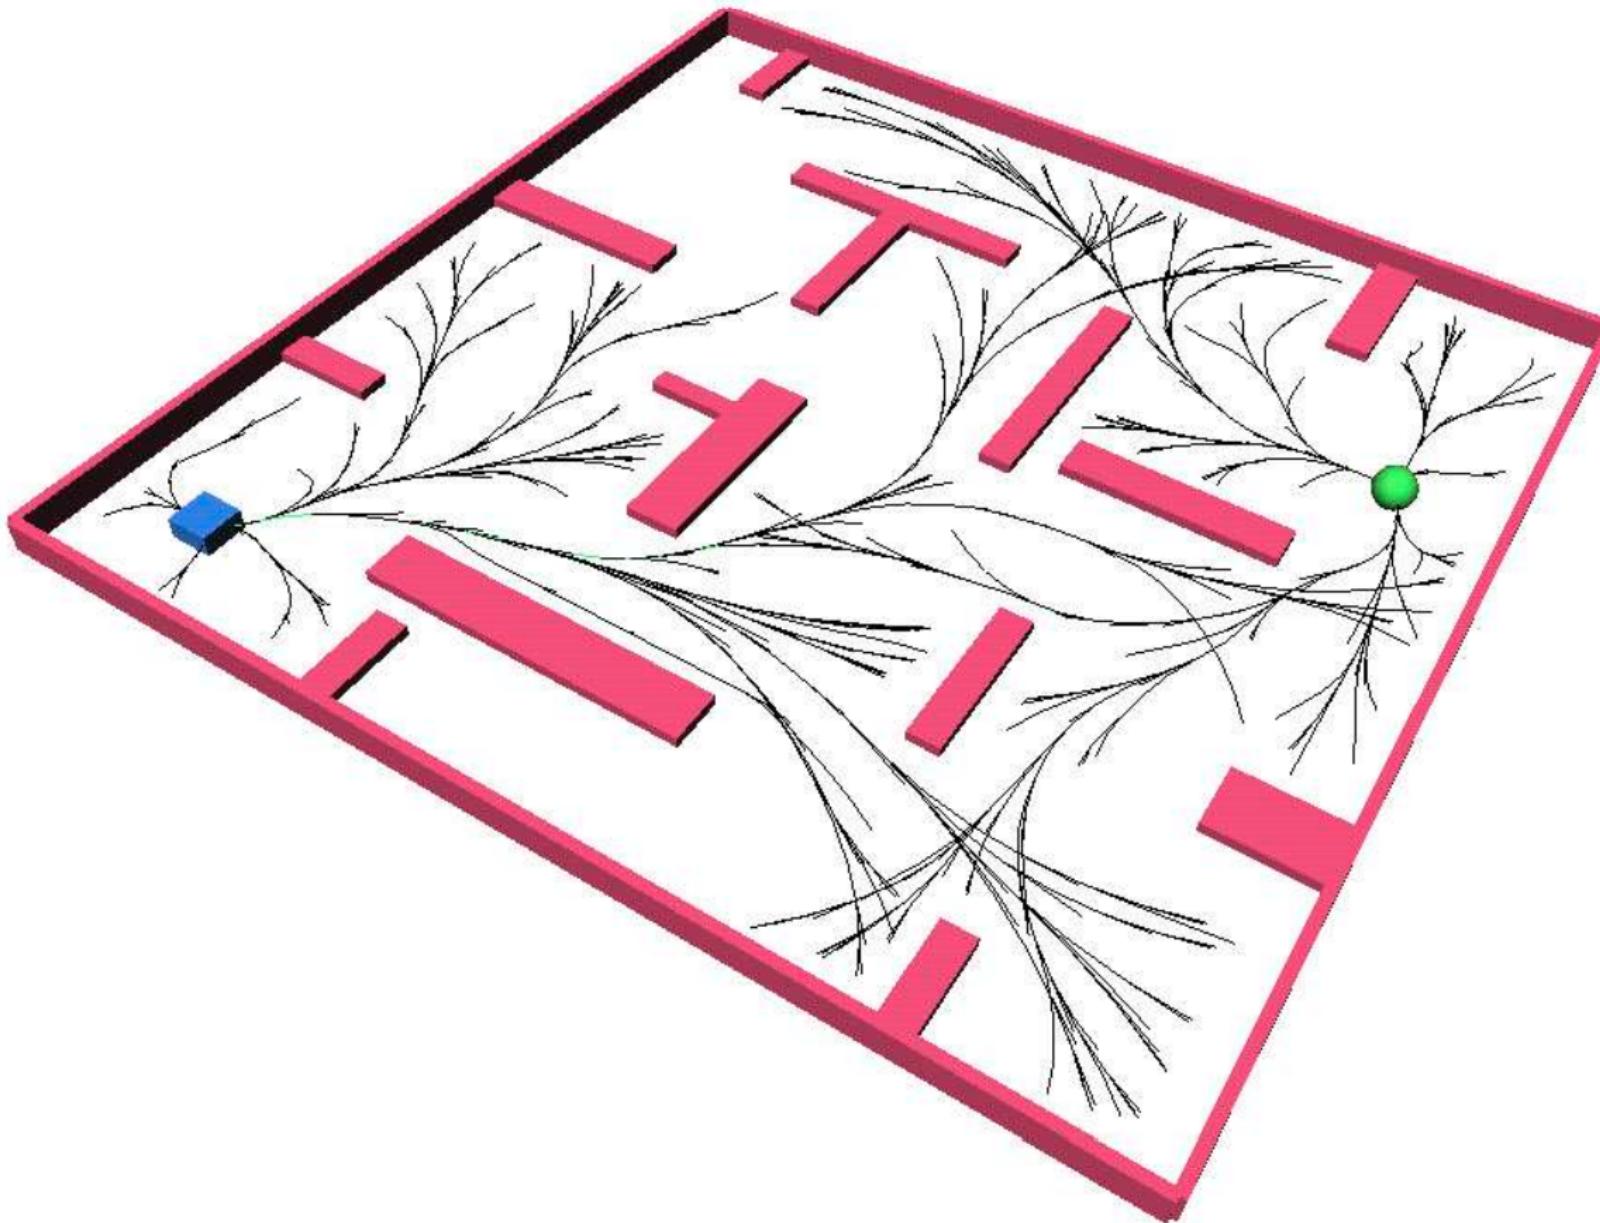
$$\mathbf{x}' = \mathbf{x} + \mathbf{f}(\mathbf{x}, \mathbf{u})$$

Without knowing anything specific about f , expanding toward a sampled state is not possible.

But RRT can be adapted, either:

- by choosing \mathbf{u} randomly (*random expansion*), or:
- by trying many possible inputs and choosing one that yields a new state as close as possible to the target sample.

Kinodynamic RRT works well too:

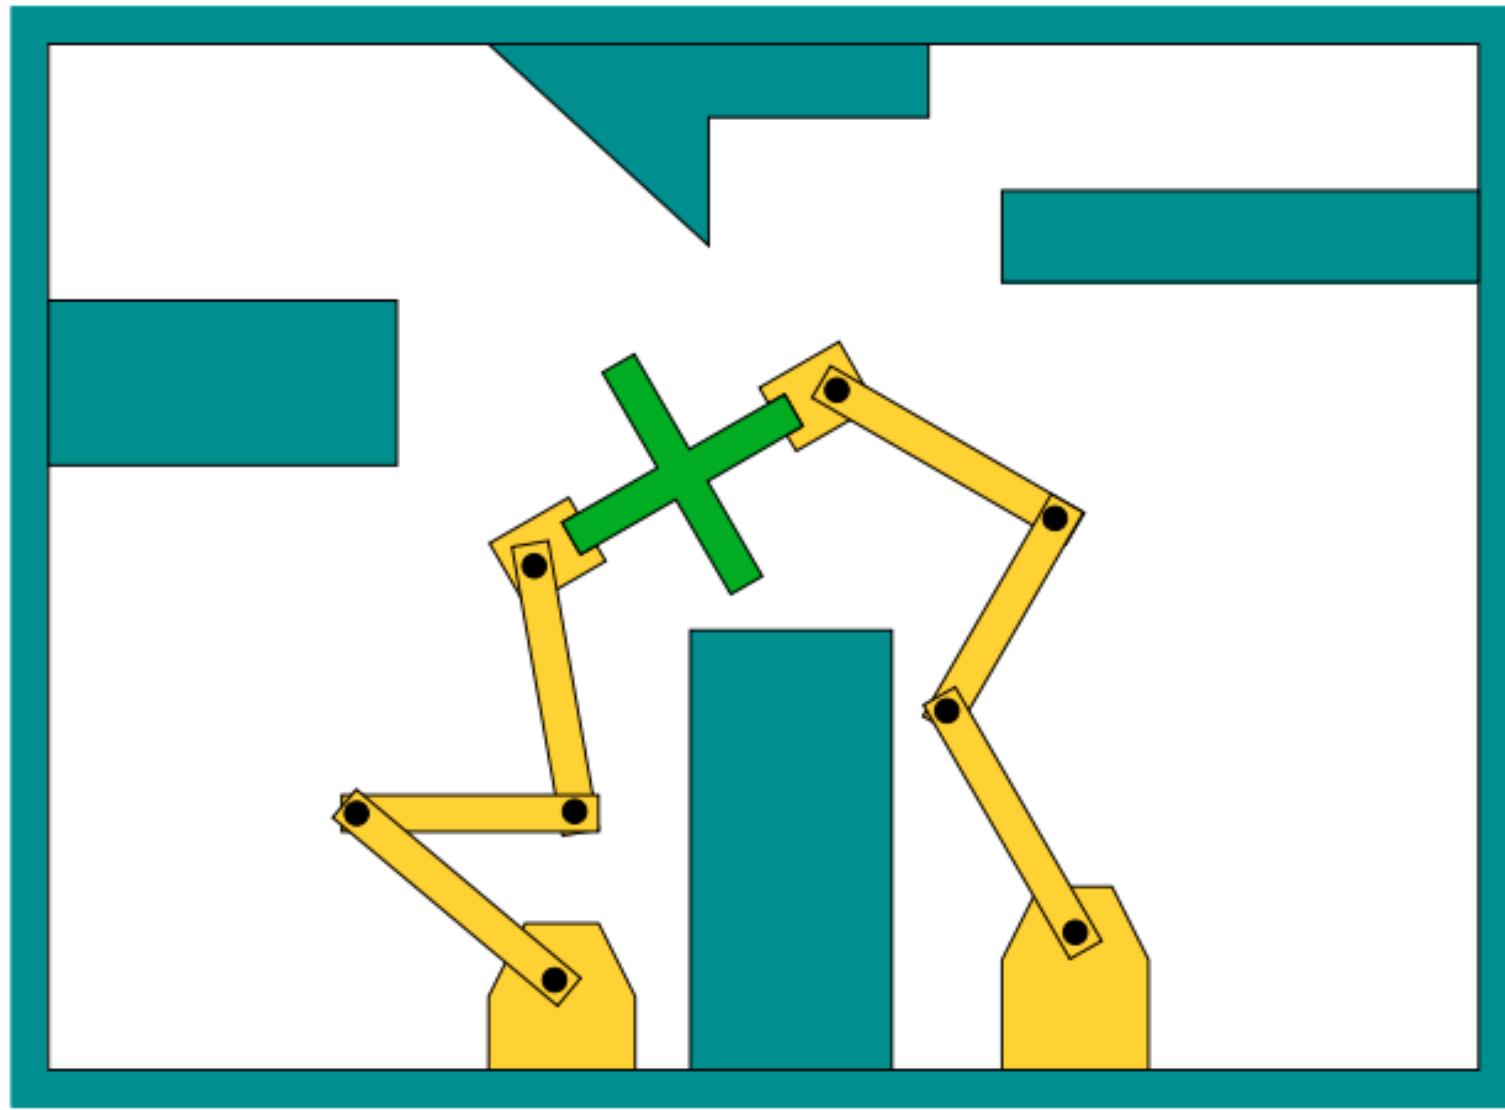


Kinodynamic RRT

Remark: techniques like RRT* cannot be applied to kinodynamic motion planning, therefore asymptotic optimality remains a big challenge for problems with dynamic constraints.

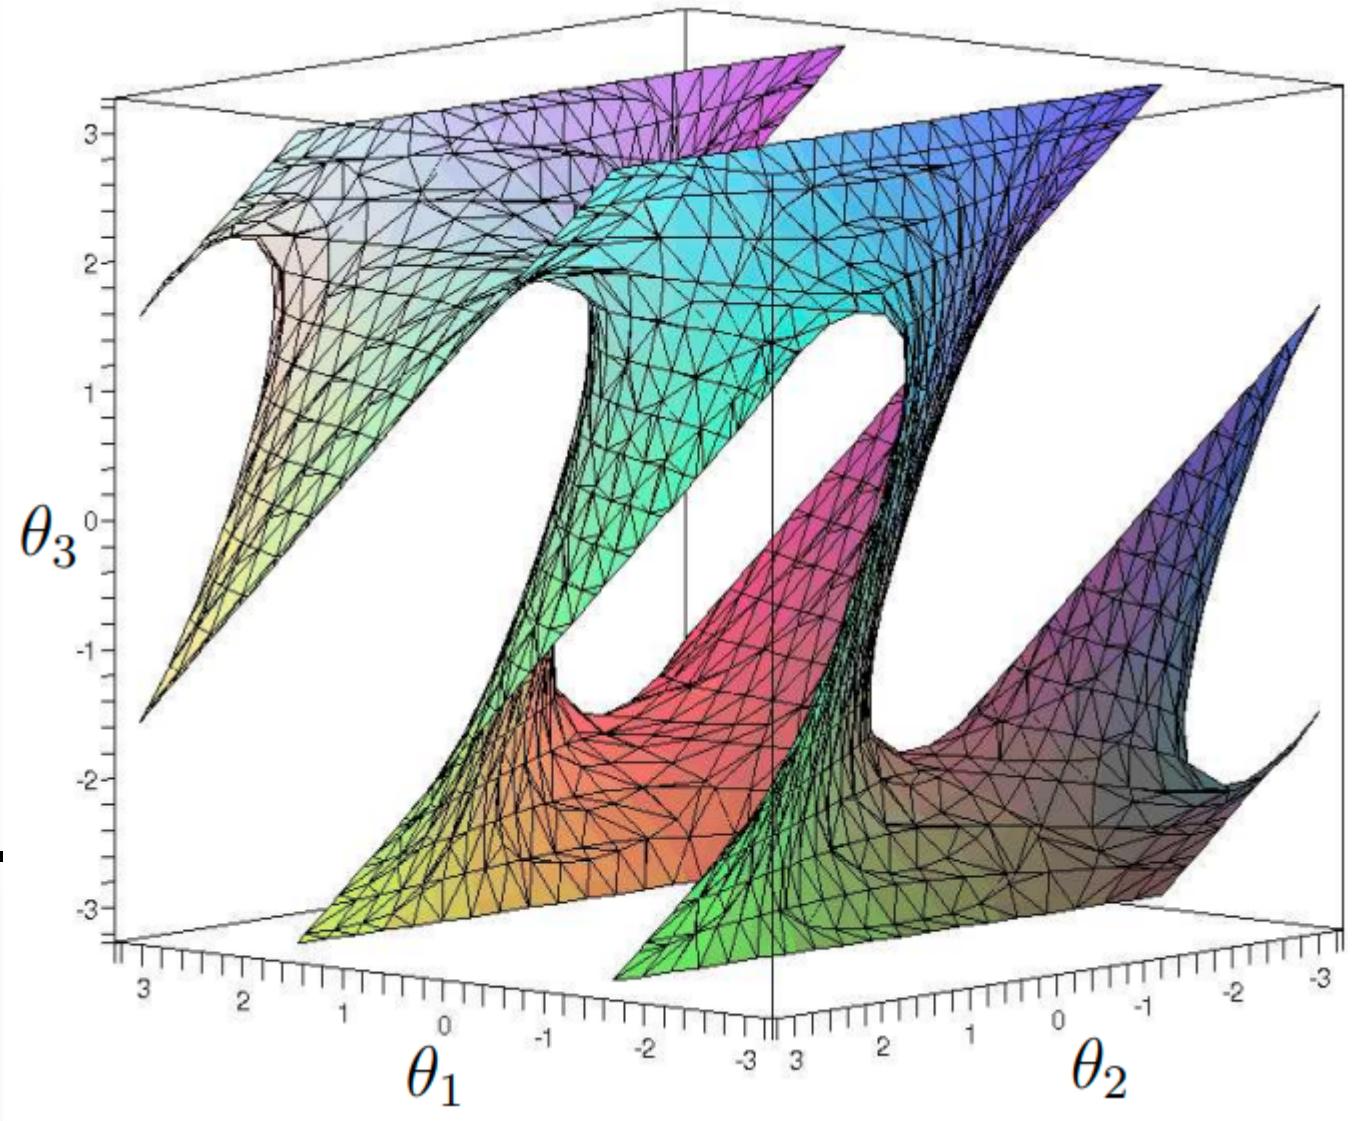
A way to address this issue is to use motion planning to find *initial trajectory candidates*, and then use another technique to incrementally improve the solution (e.g. *reinforcement learning*, *optimal control*, ...)

When it is hard to sample configurations:
example with closed kinematic chains



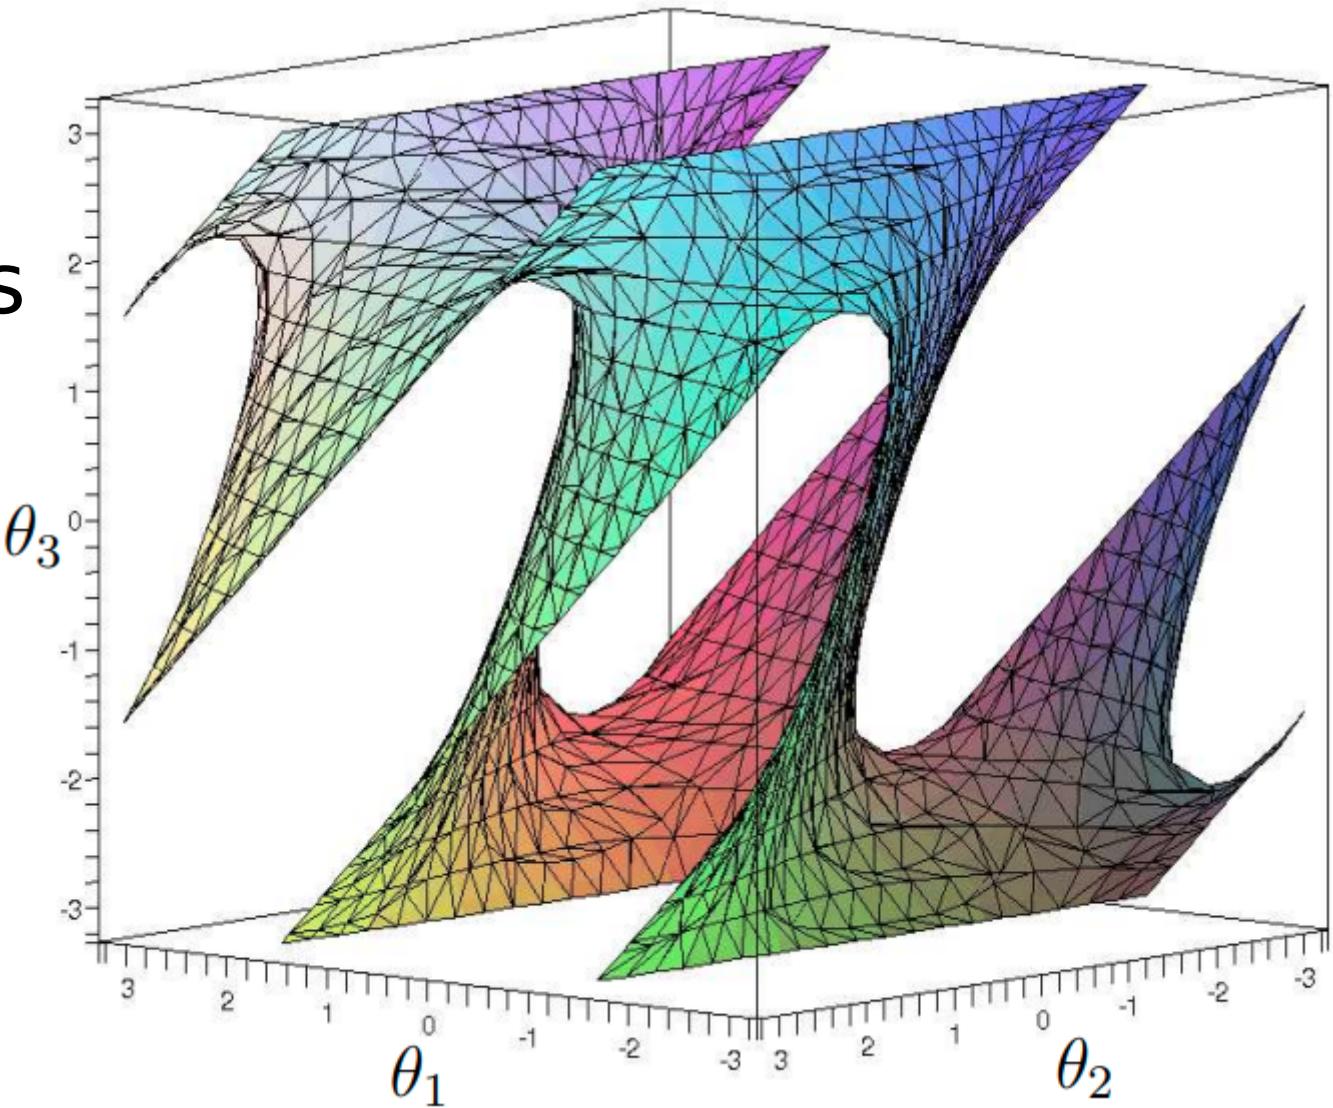
Closed kinematic chains

Example of configuration space (a submanifold of the joint space) for a 3-link kinematic chain.



Closed kinematic chains

The motion planning must explore trajectories that remain on the submanifold, but it cannot be ensured using control inputs.



Closed kinematic chains

There ways to adapt sampling-based motion planning algorithms, usually by solving loop closure equations to project all considered motions on the correct submanifold.

See: **Motion planning algorithms for general closed-chain mechanisms**, J. Cortes, 2005.

Multiple-query motion planning

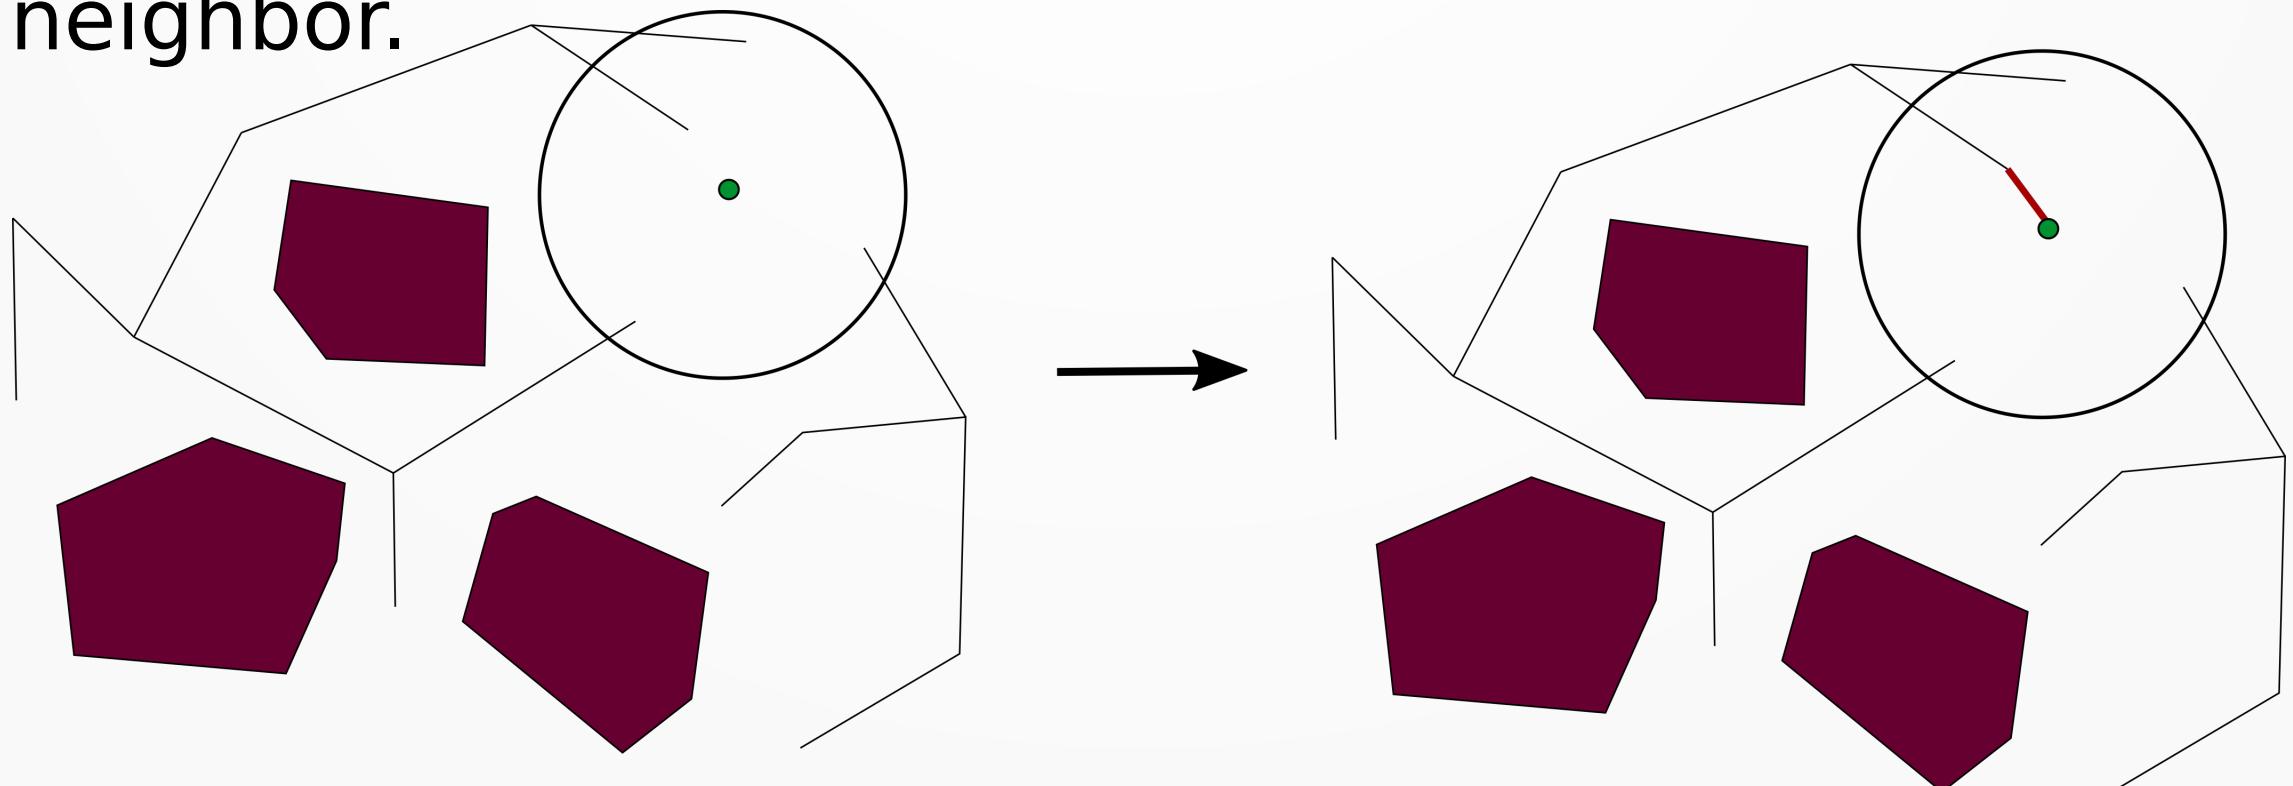
Goal: building a *roadmap*, i.e. a graph of motions that can be useful for quickly answering motion planning queries.

Probabilistic Roadmaps (PRM)

See: **Probabilistic Roadmaps for Path Planning in High-Dimensional Configuration Spaces**, L.E. Kavraki, P. Svestka, J.-C. Latombe, M. H. Overmars, IEEE Transactions on Robotics and Automation, 1996.

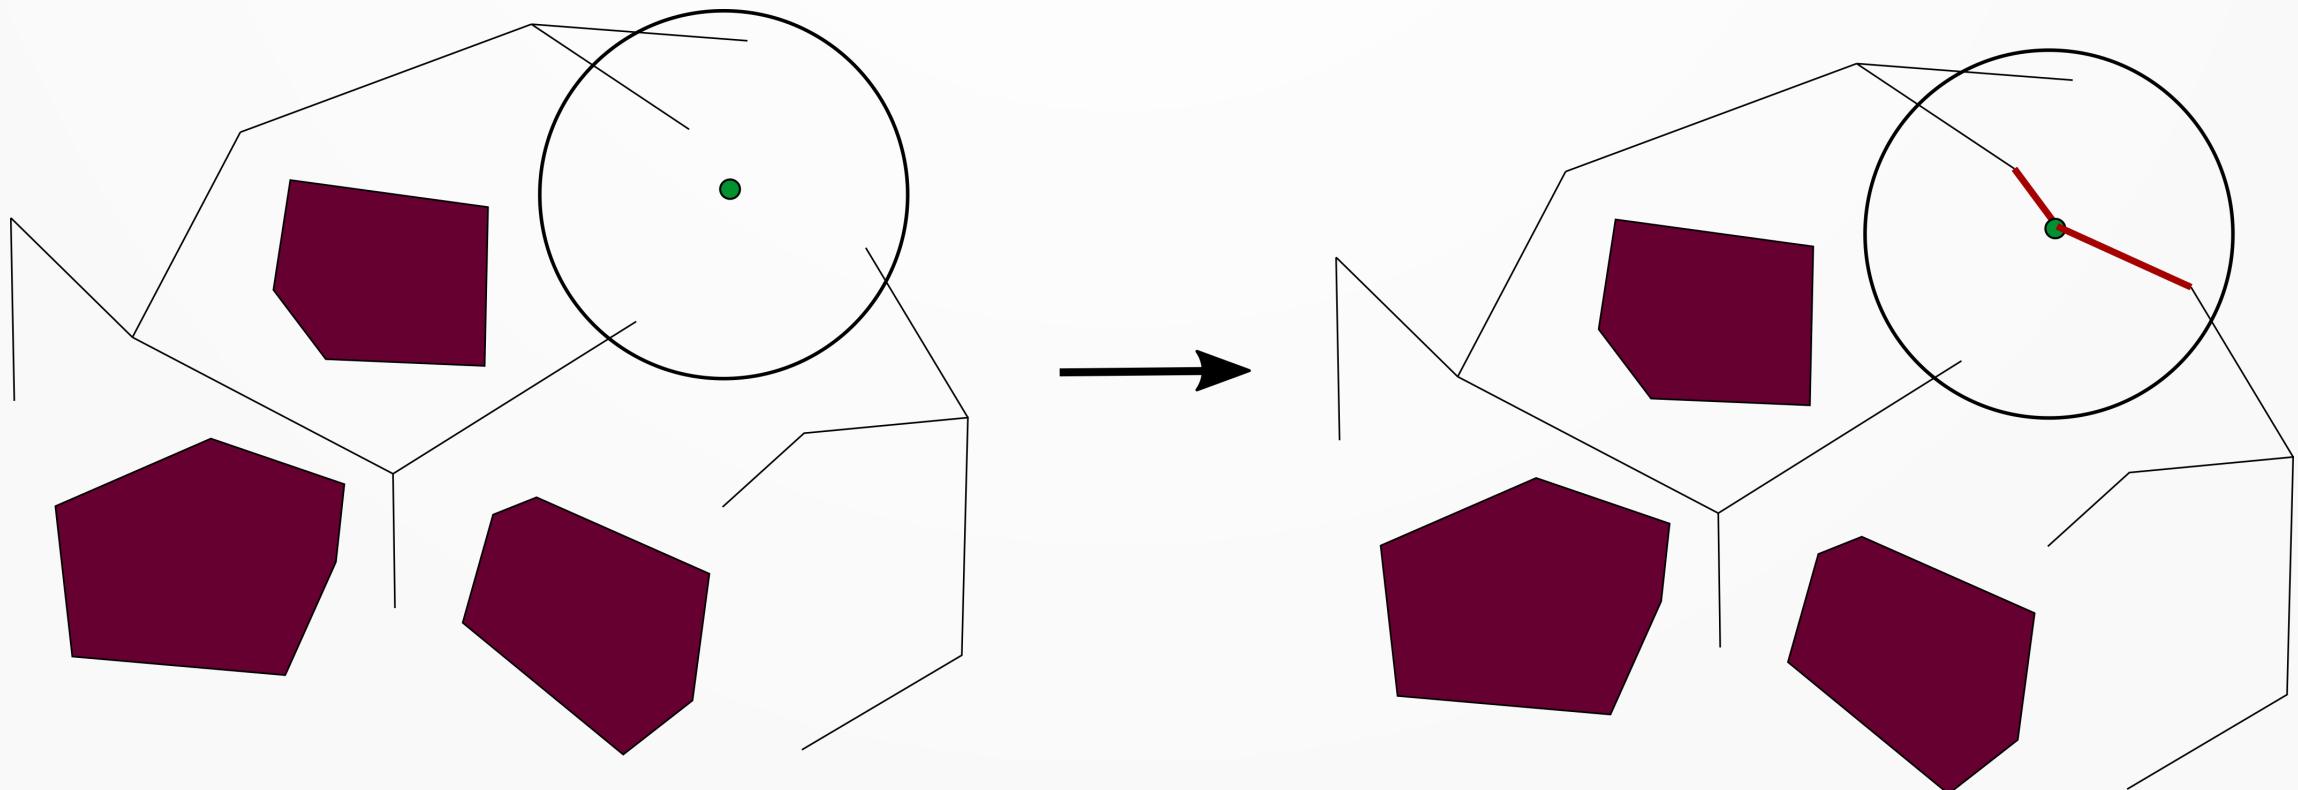
Probabilistic Roadmaps (PRM)

(1) Randomly sample a configuration, and consider neighbors in the current graph within a sphere of arbitrarily fixed radius. If possible, create a collision-free edge between the new sample and its nearest neighbor.

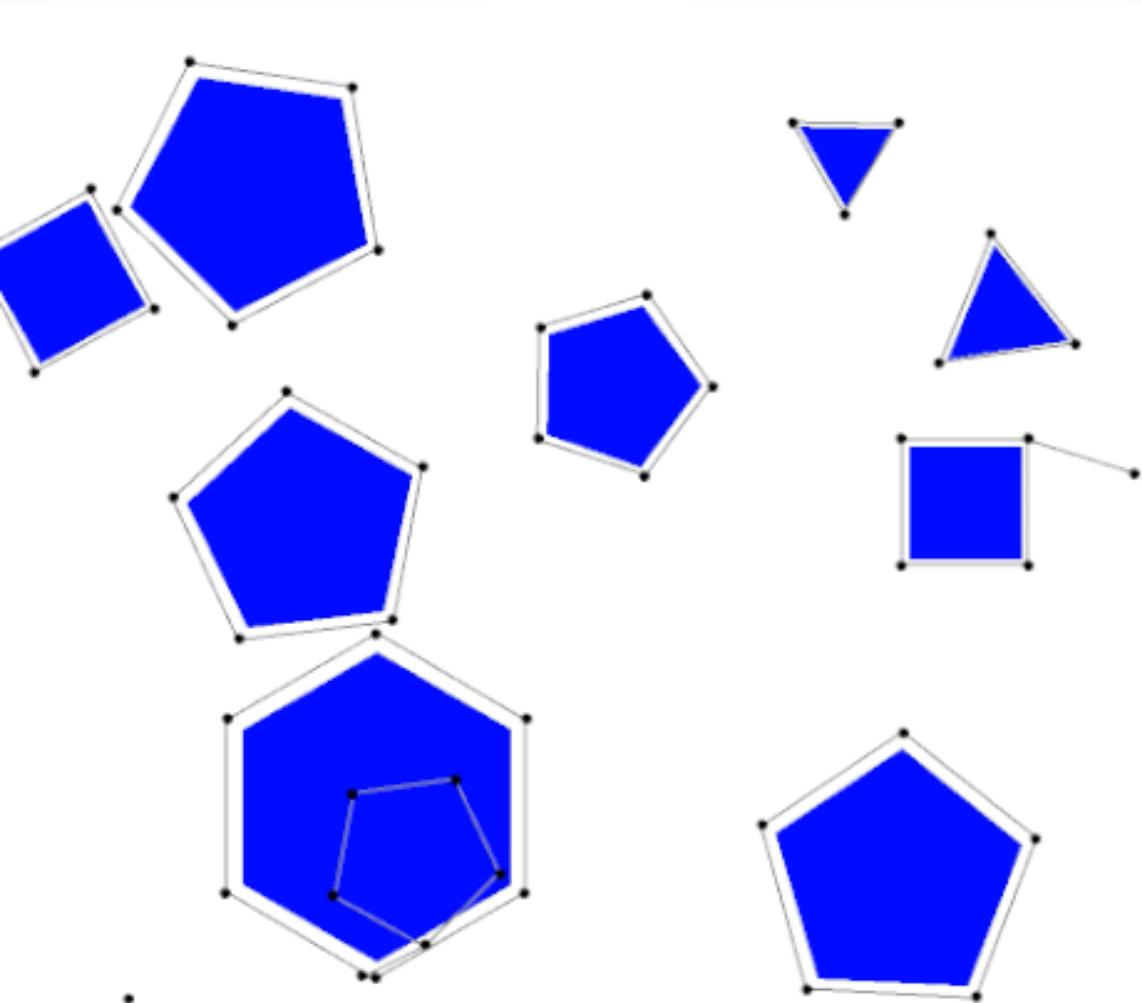


Probabilistic Roadmaps (PRM)

(2) Consider the other neighbors by order of increasing distance, and add an edge (if collision-free) if and only if it joins two distinct connected components of the graph (roadmap).



Probabilistic Roadmaps (PRM)



Probabilistic Roadmaps (PRM)

The query phase relies on a *local planner* to connect the source configuration **S** and goal configuration **G** to the roadmap, and then applies Dijkstra's algorithm to find the shortest path from **S** to **G** via the roadmap.